

Lecture: Machine Learning for Data Science

Winter semester 2021/22

Lecture 9: Regression

Prof. Dr. Eirini Ntoutsi

Outline

- Intro into regression
- Univariate linear regression
- Multivariate linear regression
- Linear classifiers with a hard/ logistic threshold
- Things you should know from this lecture & reading material

Regression vs Classification

- Both supervised learning tasks
 - In classification, the class attribute is discrete.
 - In regression, the class attribute is continuous.

ID	Age	Car type	Risk
1	23	Familie	high
2	17	Sport	high
3	43	Sport	high
4	68	Familie	low
5	32	LKW	low

What is the predicted risk of a person with a certain Age and Car type?

House ID	Size (feet)	Old	Price
1	500	10	100K
2	1000	20	500K
3	2000	50	300K
4	300	15	200K


What is the predicted price for a house of a certain size and age?

Regression example (univariate case)

- Consider the following dataset of different houses and their prices
- Given this data, a friend has a house 750 square feet -how much can they be expected to get?

instance	size	price
1	500	100
2	1000	250
3	2000	300


Housing price prediction.



Source: Andrew Ng ML course, Coursera

Regression application: Precision farming (multivariate case)

- Predict the amount of fertilizers based on multiple attributes like soil characteristics, weather, used fertilizers.
 - Only the appropriate amount of fertilizers given the environmental settings (soil, weather) will result in maximum yield.
 - Controlling the effects of over-fertilization on the environment is also important



Problem formulation

- “A computer program is said to learn from *experience* E w.r.t. some class of *tasks* T and *performance measure* P , if its performance at tasks in T , as measured by P , improves with experience E .”

[Tom Mitchell](#), Machine Learning 1997.

- What is the task T ?
- What is the experience E ?
- What is the performance/evaluation measure P ?

Problem formulation

- In **classification**, the class attribute Y is discrete: $d(Y) = \{c_1, c_2, \dots, c_k\}$; k is the number of classes.
- The goal is to find a mapping/function/hypothesis $h(): X \rightarrow Y$
- In **regression**, the prediction aims at a real value $Y \in \mathbb{R}$
- The goal is to learn a function $h()/f(): X \rightarrow \mathbb{R}$
- Each training instance has the form (\vec{x}, y) where $y \in \mathbb{R}$
- The predicted value $f(x)$ is also called **predicted variable**, response variable, or dependant variable and is often denoted by y
- Examples
 - ❑ Recommended amount of fertilizer for a certain type of soil
 - ❑ Prediction of house prices

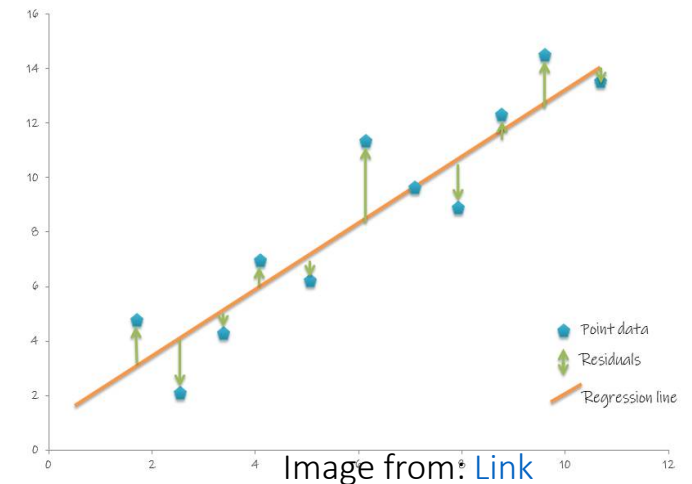
Task T

House ID	Size (feet)	Old	Price
1	500	10	100K
2	1000	20	500K
3	2000	50	300K
4	300	15	200K

Experience E

Problem formulation

- How to evaluate the performance of $h()$? Performance measure P
- Using the **empirical error** or **empirical loss** of $h()$
 - So errors over the training set
- In **classification**, we were *mainly* checking whether the predicted class agrees with the real class/ground truth for each instance, i.e., $h(x_i)=y_i$
- For **regression**, we want to check how close is the predicted value $h(x_i)$ from the real value y_i
- Popular loss functions
 - **L_1 loss or absolute-value loss**: $L_1(h(x_i), y_i) = |h(x_i) - y_i|$
 - **L_2 loss or squared-error loss**: $L_2(h(x_i), y_i) = (h(x_i) - y_i)^2$
 - **0/1 loss**: $L_{0/1}(h(x_i), y_i) = 0$, if $h(x_i) = y_i$; 1 otherwise




Outline

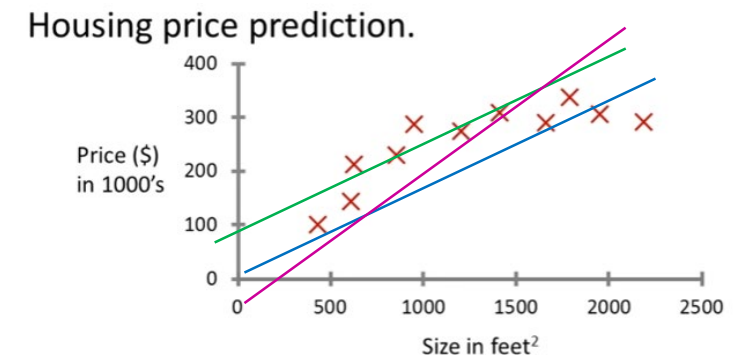
- Intro into regression
- Univariate linear regression
- Multivariate linear regression
- Linear classifiers with a hard/ logistic threshold
- Things you should know from this lecture & reading material

Univariate linear regression

$f(x)=2x+3$
offset is 3
Slope is 2



- Let a training set of N instances: $D=\{(x_i, y_i)\}$
 - each instance is described in the 1-dimensional feature space: (X)
 - The class Y is continuous
- We are looking for a univariate linear function (a straight line) with the form
$$f_w(x) = w_1 \cdot x + w_0$$
 - w_1 is the slope (orientation)
 - w_0 is the intercept (offset from origin) (**bias**)
- $w=\langle w_1, w_0 \rangle$ is the weight vector/the parameters of the model/line (to be learned from data)
- Among the available lines, which one to choose?
 - The one that best fits the data!



Univariate linear regression

- Operational definition: the one (w^*) that minimizes the empirical loss, i.e., loss over the training data
- Typically L_2 loss is used

$$Loss(h_{\mathbf{w}}) = \sum_{j=1}^N L_2(y_j, h_{\mathbf{w}}(x_j)) = \sum_{j=1}^N (y_j - h_{\mathbf{w}}(x_j))^2 = \sum_{j=1}^N (y_j - (w_1 x_j + w_0))^2$$

- So, the best line is:

$$\mathbf{w}^* = \operatorname{argmin}_{\mathbf{w}} Loss(h_{\mathbf{w}})$$




Image from: [Link](#)

How to find the best line w^*

- The goal is to find the line that minimizes the empirical L_2 error

$$\mathbf{w}^* = \operatorname{argmin}_{\mathbf{w}} \operatorname{Loss}(h_{\mathbf{w}})$$

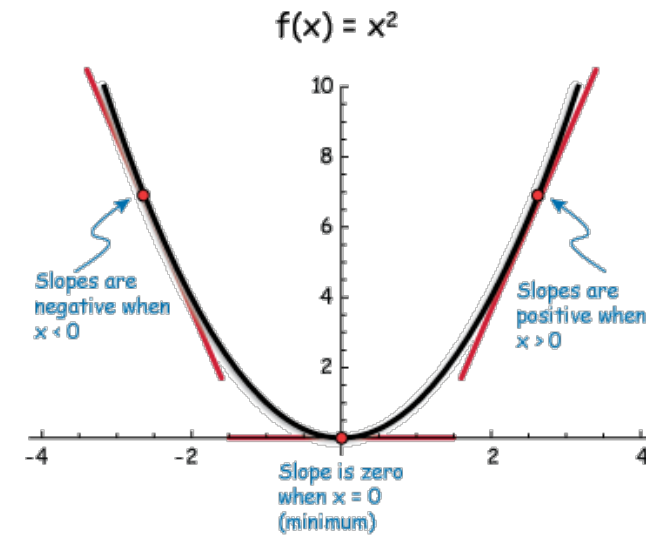
$$\operatorname{Loss}(h_{\mathbf{w}}) = \sum_{j=1}^N L_2(y_j, h_{\mathbf{w}}(x_j)) = \sum_{j=1}^N (y_j - h_{\mathbf{w}}(x_j))^2 = \sum_{j=1}^N (y_j - (w_1 x_j + w_0))^2$$

- **Analytical solution:** $\operatorname{Loss}()$ is minimized when partial derivatives w.r.t. w_1 and w_0 are zero

$$\frac{\partial}{\partial w_0} \sum_{j=1}^N (y_j - (w_1 x_j + w_0))^2 = 0 \text{ and } \frac{\partial}{\partial w_1} \sum_{j=1}^N (y_j - (w_1 x_j + w_0))^2 = 0.$$

- These equations have a unique solution:


$$w_1 = \frac{N \left(\sum x_j y_j \right) - \left(\sum x_j \right) \left(\sum y_j \right)}{N \left(\sum x_j^2 \right) - \left(\sum x_j \right)^2}; \quad w_0 = \left(\sum y_j - w_1 \left(\sum x_j \right) \right) / N.$$




How to find the best line w^*

- In this case, we can find a solution analytically because the loss function is convex
- In the general case, for any loss function, we can use **gradient-descent** (see previous lecture)
 - Start with an arbitrary initial line represented by the weight vector (w_0, w_1)
 - Repeatedly modify it in small steps
 - At each step, the weight vector is altered in the direction that produces the steepest descent along the error surface.
 - gradient descent learning rule: At each step, we take a step into the opposite direction of the gradient, and the step size is determined by the value of the **learning rate η** as well as the **slope of the gradient**
 - Recall the perceptron lecture with L2 loss

$$\Delta \vec{w} = -\eta \nabla E[\vec{w}] \quad E[\vec{w}] \equiv \frac{1}{2} \sum_{d \in D} (t_d - o_d)^2$$




Gradient Descent visualization [Link](#)




Gradient descent pseudocode

```
Initialize weight vector w
While not converged
  For each  $w_i \in w$  do
     $w_i \leftarrow w_i - \eta \frac{\partial}{\partial w_i} \text{Loss}(w)$ 
```

Univariate regression: an example



(a)



(b)

Figure 19.13 (a) Data points of price versus floor space of houses for sale in Berkeley, CA, in July 2009, along with the linear function hypothesis that minimizes squared-error loss: $y = 0.232x + 246$. (b) Plot of the loss function $\sum_j (y_j - w_1 x_j + w_0)^2$ for various values of w_0, w_1 . Note that the loss function is convex, with a single global minimum.

Source: AI book

Polynomial regression

- In the general case, the relationship between X and Y can be approximated using a larger degree polynomial
- **Polynomial regression**: A form of regression, in which the relationship between the predictive attribute X and the predictive attribute Y is modeled as a n-degree polynomial

$$f(x) = w_0 + w_1x + w_2x^2 + \dots + w_nx^n$$

- For a 1-degree polynomial → **linear regression**


$$f(x) = w_0 + w_1x$$

- For a 2-degree polynomial → quadratic

$$f(x) = w_0 + w_1x + w_2x^2$$


- For a 3-degree polynomial → cubic

$$f(x) = w_0 + w_1x + w_2x^2 + w_3x^3$$




Polynomial regression

- In many situations there exist no linear relationship between X and Y



Beware of the overfitting

- The more complex the model (higher degree n), the higher the overfitting risk




Underfitting


$$f(x) = \lambda_0 + \lambda_1 x \dots (1)$$

$$f(x) = \lambda_0 + \lambda_1 x + \lambda_2 x^2 \dots (2)$$

$$f(x) = \lambda_0 + \lambda_1 x + \lambda_2 x^2 + \lambda_3 x^3 + \lambda_4 x^4 \dots (3)$$



Best model




Overfitting

Overfitting

- Intuition: Both blue and green lines are solutions of squared loss, we prefer green as it illustrates a less complex model than the blue.
- How to avoid overfitting → by **regularizing** the complexity of the hypothesis
- **Complexity of a hypothesis $h_w()$**
 - Defined as a function of the weights w

$$\text{Complexity}(h_{\mathbf{w}}) = L_q(\mathbf{w}) = \sum_i |w_i|^q$$

- For $q=1$, minimize the sum of absolute weight values → **L_1 regularization**
 - E.g., in Lasso regression
- For $q=2$, minimize the sum of squared values → **L_2 regularization**
 - E.g., in Ridge regression
- The goal then is to find the hypothesis h^* , with the minimum total cost
$$\text{Cost}(h) = \text{Loss}(h) + \lambda \text{Complexity}(h)$$
- By adding a penalty to the loss function the overfitting is reduced



Source:
[https://en.wikipedia.org/wiki/Regularization_\(mathematics\)](https://en.wikipedia.org/wiki/Regularization_(mathematics))

L2 vs L1 regularization

- Which regularization to choose depends on the problem at hand
- L_1 regularization tends to produce **sparse models**
 - Larger penalties result in weights closer to zero, which is the ideal for producing simpler models
 - In practice, this means setting many feature weights to zero
 - Effectively, the corresponding features do not count for the prediction
- Important property
 - for interpretability
 - Less likely for the model to overfit

Outline

- Intro into regression
- Univariate linear regression
- Multivariate linear regression
- Linear classifiers with a hard/ logistic threshold
- Things you should know from this lecture & reading material

Multivariate linear regression

- Let a training set of N instances: $D=\{(\vec{x}_i, y_i)\}$
 - each instance is described in the n -dimensional feature space: (X_1, X_2, \dots, X_n)
 - The class Y is continuous
- We are looking for a **multivariate linear function** (not a line anymore but a hyperplane) with the form


$$f_w(x) = w_1 \cdot x_1 + w_2 x_2 + \dots + w_n x_n + w_0 = \sum_{i=0}^n w_i x_i$$

- Weight vector w consists of: the weights of the features w_1, \dots, w_n and the bias term w_0
- The weight vector are the parameters of the model/line (to be learned from data)
- Among the available models, which one to choose?
 - The one that best fits the data \rightarrow minimize L_2 error over the training data

$$\mathbf{w}^* = \underset{\mathbf{w}}{\operatorname{argmin}} \sum_j L_2(y_j, \mathbf{w} \cdot \mathbf{x}_j)$$


- Can be solved analytically or via gradient descent

Regression line vs regression (hyper)plane




[Link](#)

Regression line/ regression (hyper)plane vs classification line/(hyper)plane




[Link](#)

A hyperplane in \mathbb{R}^2 is a line



A hyperplane in \mathbb{R}^3 is a plane



A hyperplane in \mathbb{R}^n is an $n-1$ dimensional subspace

Overfitting (discussed already)

- In multivariate regression, some attribute/dimension that is irrelevant might by chance appear to be useful → **overfitting**
- How to avoid overfitting → by **regularizing** the complexity of the hypothesis (already discussed)
- **Complexity of a hypothesis $h_w()$**
 - Defined as a function of the weights w

$$Complexity(h_{\mathbf{w}}) = L_q(\mathbf{w}) = \sum_i |w_i|^q$$


- For $q=1$, minimize the sum of absolute weight values → **L_1 regularization**
 - For $q=2$, minimize the sum of squared values → **L_2 regularization**
- The goal then is to find the hypothesis h^* , with the minimum total cost
$$Cost(h) = Loss(h) + \lambda Complexity(h)$$

Outline

- Intro into regression
- Univariate linear regression
- Multivariate linear regression
- Linear classifiers with a hard/ logistic threshold
- Things you should know from this lecture & reading material


Linear classifiers with a hard threshold

- We have already seen that linear functions can be also used for classification
 - The decision boundary is a line or hyperplane in higher dimensional spaces that separates the classes
- Recall the **simple perceptron classifier**



Step function
(threshold=0)

Activation function is the step function



$$o(\vec{x}) = \begin{cases} 1 & \text{if } \vec{w} \cdot \vec{x} > 0 \\ -1 & \text{otherwise.} \end{cases}$$


- In case of correct predictions, weights do not change
- In case of wrong predictions, weights change following the update rule
 - $W_j \leftarrow w_j + \eta(y^i - o^i)x_j^i$

0/1 loss:


$L_{0/1}(h(x_i), y_i) = 0$, if $h(x_i) = y_i$; 1 otherwise

Convergence problems

- Although the perceptron learning rule converges if the two classes can be separated by linear hyperplane, problems arise if the classes cannot be separated perfectly by a linear classifier.




Plot of two seismic data parameters, **body wave magnitude (x_1)** and **surface wave magnitude (x_2)** for **earthquakes (open orange circles)** and **nuclear explosions (green circles)** occurring between 1982 and 1990 in Asia and the Middle East (Kebeasy *et al.*, 1998). The earthquakes and explosions are **not linearly separable**.




The simple perceptron learning rule fails to converge even after 10000 steps

From hard thresholds to soft thresholds

- The hard nature of the threshold of the simple perceptron causes problems with convergence
 - The hypothesis $h_w(x)$ is not differentiable
 - The classifier makes completely confident class predictions, even for instances close to the boundary
- Such problems can be solved using a **soft threshold approach**, so approximating the hard threshold with a continuous, differentiable function
- A popular choice is the **logistic function**, also known as **sigmoid function**



Hard- threshold/ Step
function with 0/1 output




Logistic function

$$\sigma(z) = \frac{1}{1 + e^{-z}}$$

Sigmoid unit


- Very much like the perceptron unit but based on a smoothed, differentiable threshold function




- The sigmoid output $\sigma()$ is a continuous function of its input
 - takes real values between -1 and +1
 - The sigmoid is in effect an approximation to the threshold function above, but has a gradient that we can use for learning

$$\frac{d\sigma(x)}{dx} = \sigma(x)(1 - \sigma(x))$$

An example



Plot of two seismic data parameters, **body wave magnitude (x_1)** and **surface wave magnitude (x_2)** for **earthquakes (open orange circles)** and **nuclear explosions (green circles)** occurring between 1982 and 1990 in Asia and the Middle East (Kebeasy *et al.*, 1998). The earthquakes and explosions are **not linearly separable**.




A logistic regression hypothesis.
The output is continuous in the $[0-1]$ range


The output can be interpreted as *probability* of belonging to the class labeled 1.

Logistic regression


- The process of fitting the weights of this model to minimize loss on a data set is called **logistic regression**
- The weights are computed using gradient descent (with L2 loss function)
 - (Weight update formula not shown)
- Convergence example



Plot of two seismic data parameters, **body wave magnitude (x_1)** and **surface wave magnitude (x_2)** for **earthquakes (open orange circles)** and **nuclear explosions (green circles)** occurring between 1982 and 1990 in Asia and the Middle East (Kebeasy *et al.*, 1998). The earthquakes and explosions are **not linearly separable**.



The simple perceptron learning rule fails to converge even after 10000 steps



Logistic regression converges fast

Outline

- Intro into regression
- Univariate linear regression
- Multivariate linear regression
- Linear classifiers with a hard/ logistic threshold
- Things you should know from this lecture & reading material

Overview and Reading

- Overview

- Univariate linear regression
- Multivariate linear regression
- Hard threshold linear classifiers
- Logistic regression

- Reading

- Artificial Intelligence, A Modern Approach. Stuart Russell and Peter Norvig (Chapter 19)
- Chapter 9: Linear predictors, Understanding Machine Learning book by Shai Shalev-Schwartz and Shai Ben-David

Hands on experience

- Try regression on the crime-prediction dataset (128 attributes) [Link](#)
- Many more datasets available on [UCI for regression tasks](#)



Acknowledgements

- The slides are based on
 - ❑ KDD I lecture at LMU Munich (Johannes Aßfalg, Christian Böhm, Karsten Borgwardt, Martin Ester, Eshref Januzaj, Karin Kailing, Peer Kröger, Eirini Ntoutsi, Jörg Sander, Matthias Schubert, Arthur Zimek, Andreas Züfle)
 - ❑ Introduction to Data Mining book slides at <http://www-users.cs.umn.edu/~kumar/dmbook/>
 - ❑ Pedro Domingos Machine Lecture course slides at the University of Washington
 - ❑ Machine Learning book by T. Mitchel slides at <http://www.cs.cmu.edu/~tom/mlbook-chapter-slides.html>
 - ❑ C. Burges, [A Tutorial on Support Vector Machines for Pattern Recognition](#), DMKD 1998
 - ❑ For regression, images come mainly from Artificial Intelligence, A Modern Approach. Stuart Russell and Peter Norvig (Chapter 19) book
 - ❑ Thank you to all TAs contributing to their improvement, namely Vasileios Iosifidis, Damianos Melidis, Tai Le Quy, Han Tran.

Thank you

Questions/Feedback/Wishes?