

Lecture: Machine Learning for Data Science

Winter semester 2021/22

Lecture 16: Reinforcement Learning (Intro and MDPs)

Prof. Dr. Eirini Ntoutsi

Outline

- Introduction
- MDP formulation
- Solving the MDP
- From value-iteration to policy-iteration
- Relaxing the (full) MDP assumptions → RL
- Model-based learning
- Things you should know from this lecture & reading material

Main machine learning tasks

- Based on the feedback we have on the data, we can distinguish between:

- **Direct-feedback** instances

Supervised learning

- the correct response /label is provided for each instance by the “teacher”
- e.g., good or bad product

- **No-feedback** instances

Unsupervised learning

- no evaluation/label of the instance is provided, since there is no “teacher”
- e.g., no information on whether a product is good or bad, just the description of the product/instance

- **Indirect-feedback** instances

Reinforcement learning

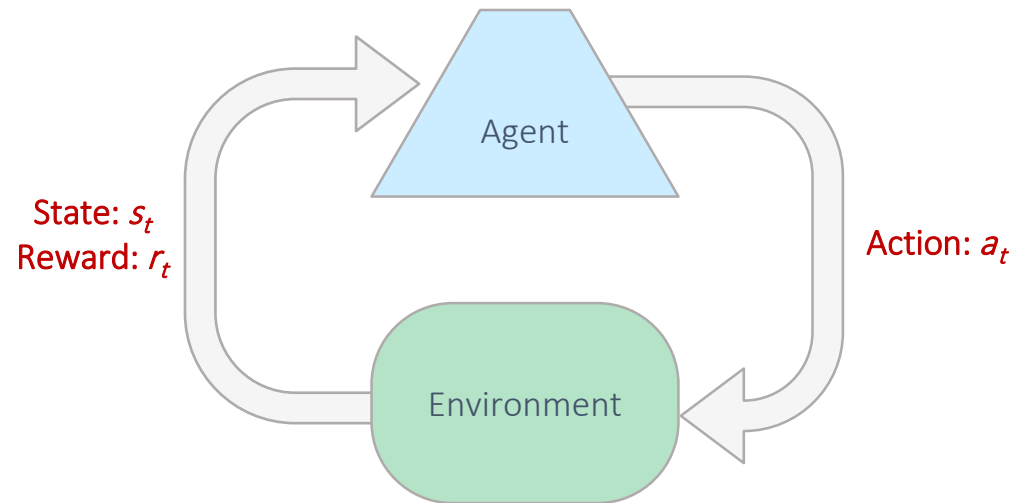
- less feedback is given, since not the proper action, but only an evaluation of the chosen action is given by the “teacher”

fruit	length	width	weight	label
fruit 1	165	38	172	Banana
fruit 2	218	39	230	Banana
fruit 3	76	80	145	Orange
fruit 4	145	35	150	Banana
fruit 5	90	88	160	Orange
...				
fruit n

fruit	length	width	weight
fruit 1	165	38	172
fruit 2	218	39	230
fruit 3	76	80	145
fruit 4	145	35	150
fruit 5	90	88	160
...			
fruit n

Agent and environment

- RL is a type of ML technique that enables an **agent** to learn in an interactive **environment** by trial and error using **feedback** from its own **actions** and **experiences**.




In the terminology of psychology,
reward is reinforcement

- At each step t
 - The agent executes action a_t
 - Transitions to state s_t
 - Receives scalar reward r_t
 - reward can be sparse
- The goal of the agent is to learn to choose actions so as to maximize the sum of rewards

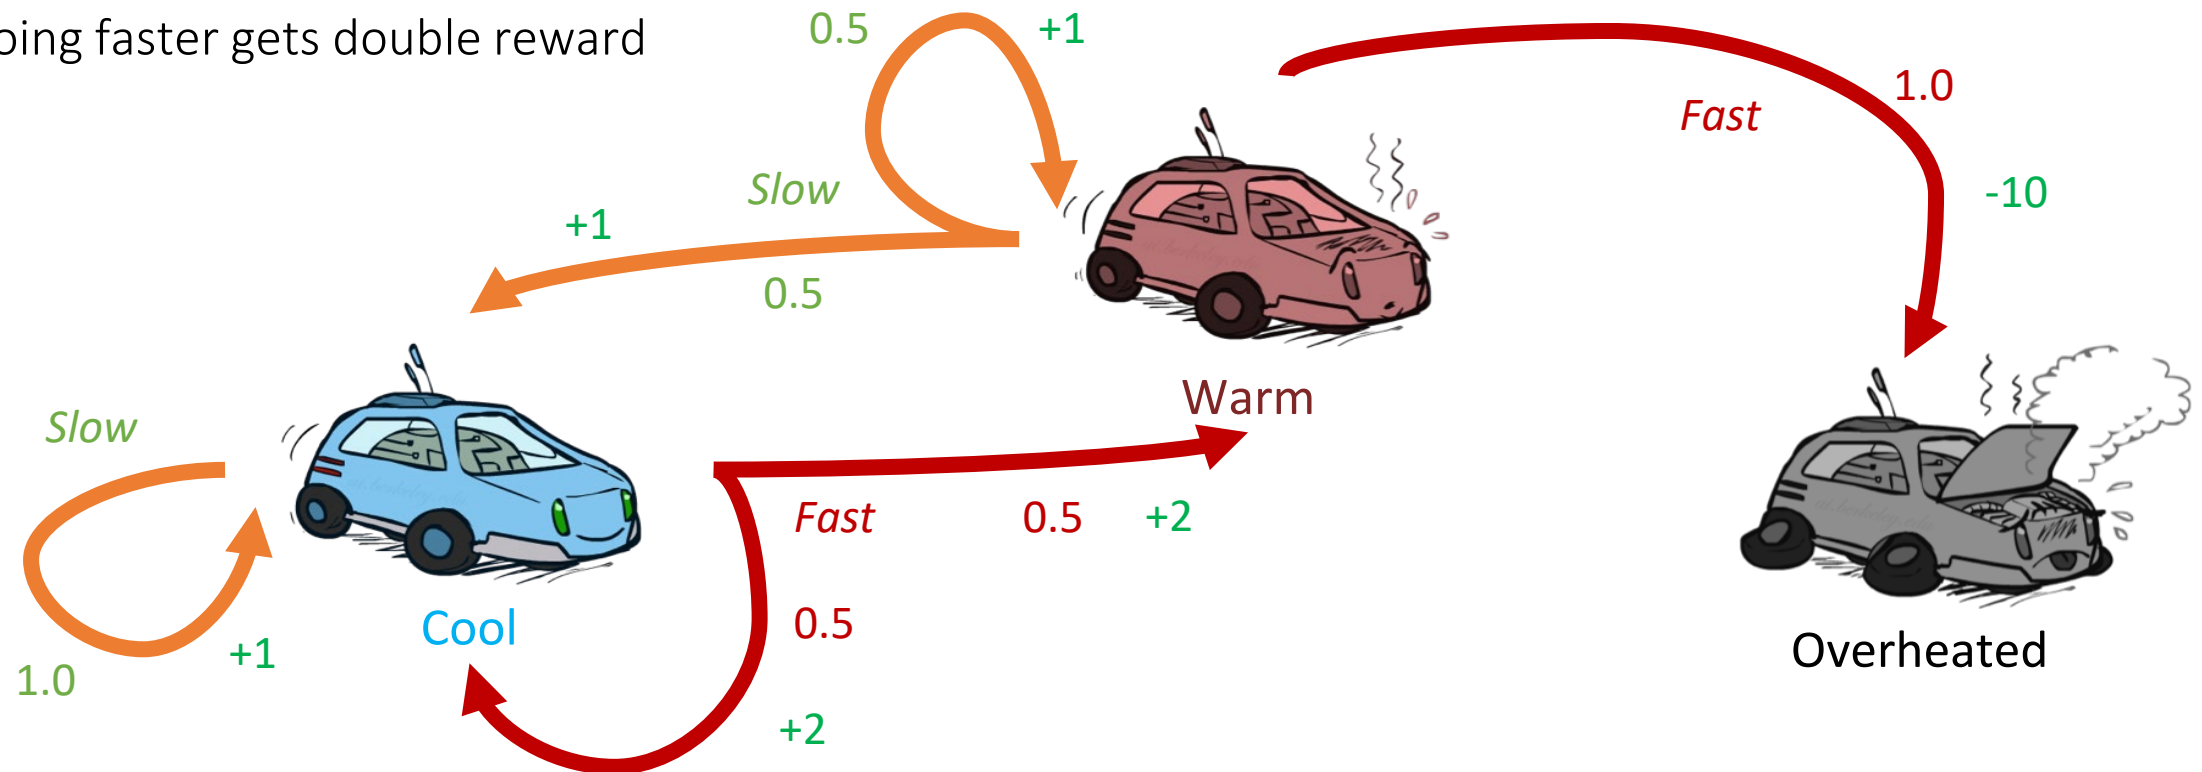
Example: Grid World

- A maze-like problem
 - The agent lives in a grid, walls block its path
 - States: different positions of the agent in the grid, assuming that wall, diamond, fire do not move
 - Actions: {North, West, East, South}
- The agent receives **rewards at each time step**
 - Indicates how well agent is doing at step t
 - Small “living” reward each step (can be negative), e.g., -0.04
 - Big rewards come at the end (good or bad)
- **Uncertainty**: actions do not always go as planned
 - E.g., 80% of the time, the action North takes the agent North (if there is no wall there)
 - 10% of the time, North takes the agent West; 10% East
 - If there is a wall in the direction the agent would have been taken, the agent stays put
- **Goal**: maximize sum of rewards



Example: Racing

- A robot car wants to travel far, quickly
- Three states: Cool, Warm, Overheated (terminal state)
- Two actions: Slow, Fast
- Going faster gets double reward




Example applications


- Play Backgammon
 - +/-ve reward for winning/losing a game
- Manage an investment portfolio
 - +ve reward for each \$ in bank
- Make a robot walk
 - +ve reward for forward motion
 - -ve reward for falling over
- Taxi driving
 - +ve reward for getting closer
 - -ve reward for moving away from the destination
- ...

Characteristics of RL: non-deterministic decision making

- **Non-deterministic**: There is uncertainty associated with the actions
 - Actions might result in multiple successor states

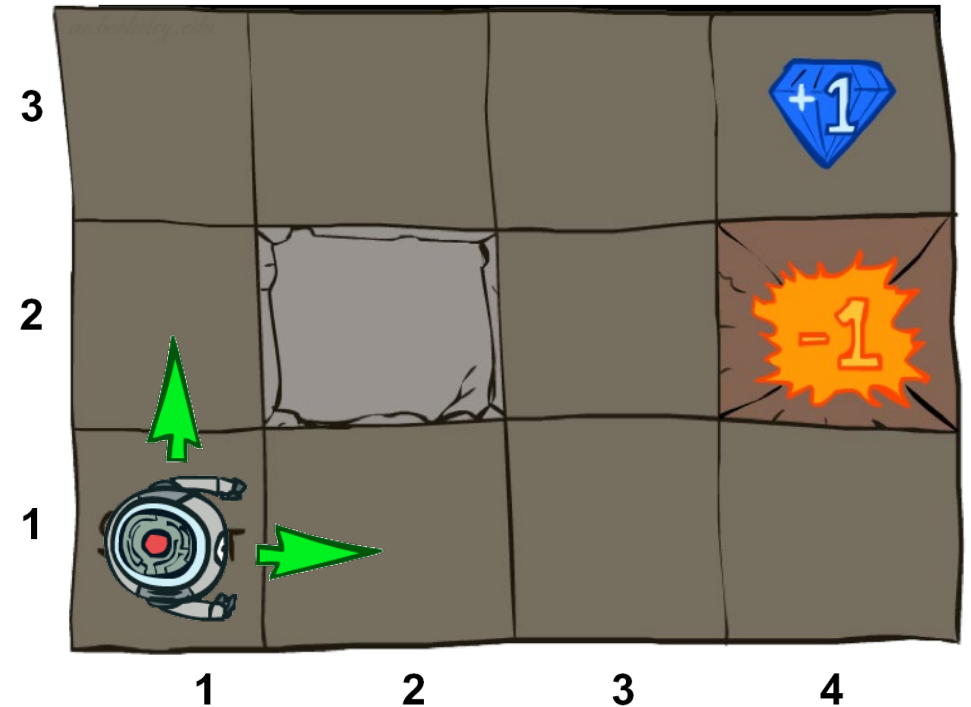


Deterministic grid world



Characteristics of RL: sequential decision making

- **Sequential decision making:**
 - the current decision could affect all future decisions.
 - short-term actions can have long-term consequences.
- Remember the goal of the agent is to maximize the sum of rewards
 - It may be better to sacrifice immediate reward to gain more long-term reward
- Examples:
 - A financial investment (may take months to mature)
 - Refuelling a helicopter (might prevent a crash in several hours)
 - Blocking opponent moves (might help winning chances many moves from now)



What makes RL different from other ML tasks?

- Agent receives feedback in the form of **rewards**
 - There is no supervisor/teacher, only a **reward** signal
- Feedback is delayed, not instantaneous
- Time really matters (sequential, non i.i.d data)
- Agent's actions affect the subsequent data it receives

How the agent learns?

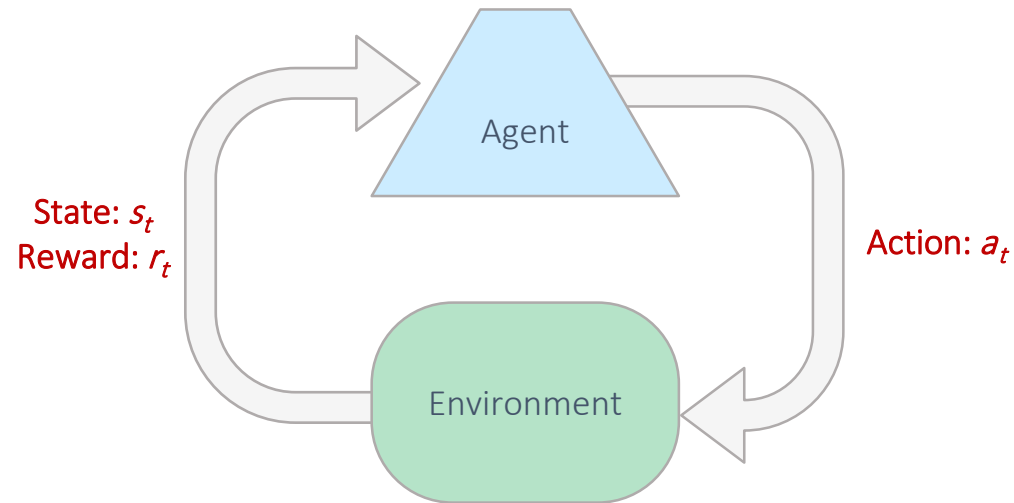
- The goal of the agent is to learn to choose actions so as to maximize the sum of rewards
- How the agent **learns**?
 - By trying out actions and observing the outcomes ← data
 - More on this later

Outline

- Introduction
- MDP formulation
- Solving the MDP
- From value-iteration to policy-iteration
- Relaxing the (full) MDP assumptions → RL
- Model-based learning
- Things you should know from this lecture & reading material

Markov Decision Process (MDP) formulation

- The interaction of the agent with the environment is modeled as a Markov Decision Process (MDP)



Markov Decision Problem (MDP) formulation

- A (full) MDP is defined by:

- A set of **states** $s \in S$

- S_0 the start/initial state
- Maybe a terminal state

- A set of **actions** $a \in A$

- $Actions(s)$: available actions in state s

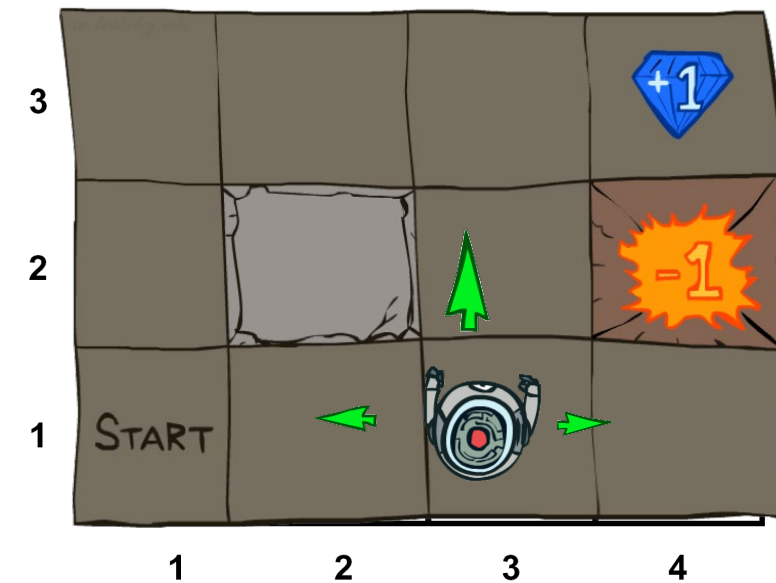
- A **transition model** $T(s, a, s')$: $T(s, a, s')$ is the probability that state s' is reached, if action a is executed in state s .

- Sometimes also, as $P(s' | s, a)$
- Transitions are **Markovian**

- **Reward function** $R(s, a, s')$: At each step the agent receives a reward

- Small living reward, higher reward (bad, good) at the terminal states
- Sometimes just $R(s)$ or $R(s')$

For the moment, we assume a full model of the world - We will relax this definition later



What is Markov about MDPs?

- “Markov” generally means “The future is independent of the past given the present”
- For Markov decision processes (MDPs), “Markov” means action outcomes depend only on the current state

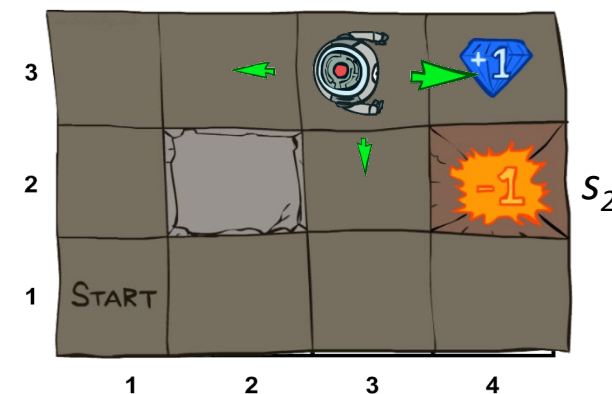
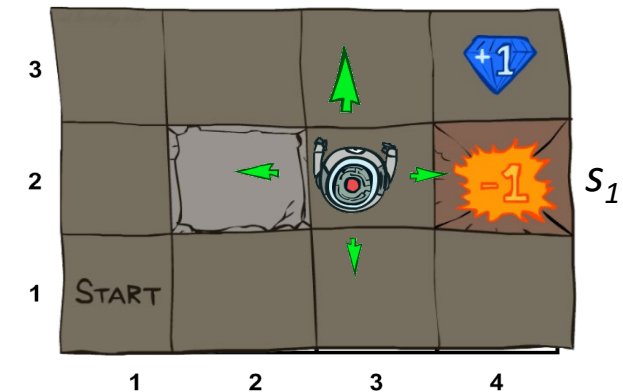
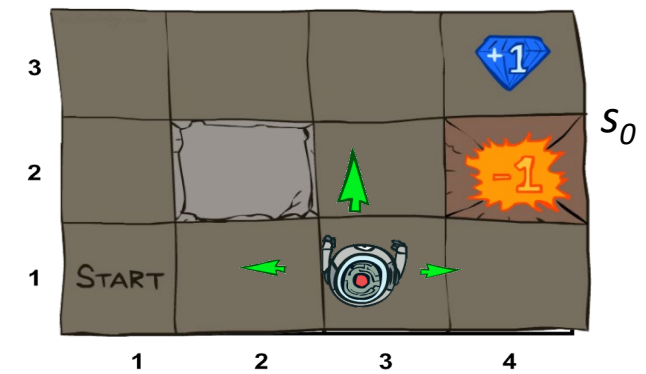
$$\begin{aligned} &P(S_{t+1} = s' | S_t = s_t, A_t = a_t, S_{t-1} = s_{t-1}, A_{t-1}, \dots, S_0 = s_0) \\ &= \\ &P(S_{t+1} = s' | S_t = s_t, A_t = a_t) \end{aligned}$$



Andrey Markov
(1856-1922)

Environment history and utility of the agent

- **Environment history**: the sequence of states $[s_0, s_1, \dots, s_n]$ “experienced” by the agent
- The **utility** of the agent depends on the environment history!
 - **Utility = sum of rewards**
- The term utility comes from AI (utility-based agents act based not only goals but also the best way to achieve the goal)
 - Utility-based agents help to choose the best alternatives, when there are multiple alternatives available.

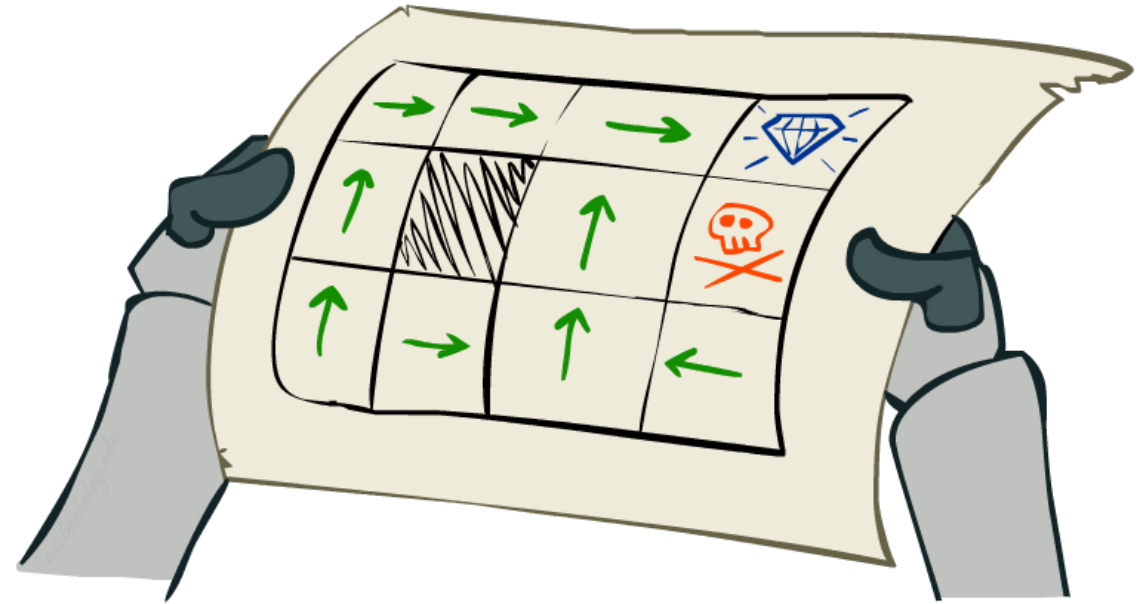



What sort of solutions we are looking for?

- Given the many choices there are many possible solutions
- A solution must specify what the agent should do for any state that the agent might reach → policy
- A policy $\pi: S \rightarrow A$ defines the agent's behavior for each state $\pi(s)$
 - Deterministic policy: $\pi(s)=a$
 - Non-deterministic/Stochastic policy: $\pi(s)$ is a distribution over possible actions given s_i
- How can we evaluate a policy π ?
 - Simple answer: By checking its utility (=sum of rewards)
 - But each time a policy π is executed starting from S_0 , the stochastic nature of the environment might lead to a different environment history generated by the policy:
$$[s_0, s_1, \dots, s_n], [s_0, s_2, \dots, s_m], \dots, [s_0, s_4, \dots, s_n]$$
 - So, to evaluate π , we need to measure the expected utility of the possible environment histories generated by that policy

Optimal policy π^*


- Optimal policy $\pi^*: S \rightarrow A$
 - An optimal policy is a policy that yields the highest expected utility



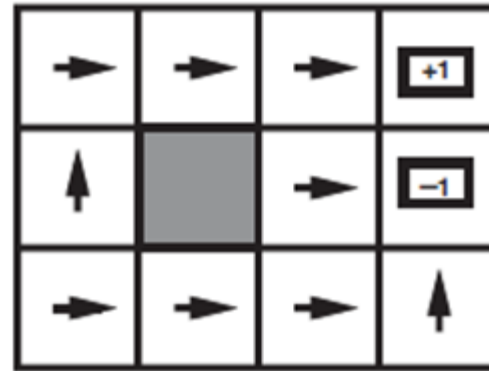
Optimal policy when $R(s, a, s') = -0.03$ for all non-terminals s

Optimal policies: balancing risk and reward


- Optimal policy changes with choice of living rewards $R(s)$.




$$R(s) = -0.04$$



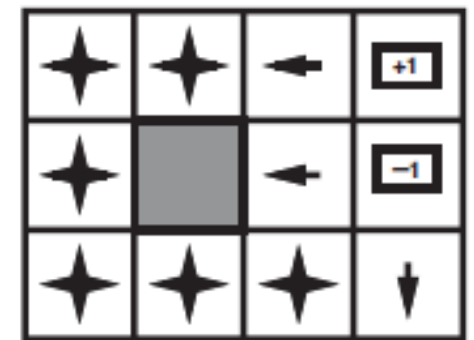
$$R(s) < -1.6284$$



$$-0.4278 < R(s) < -0.0850$$



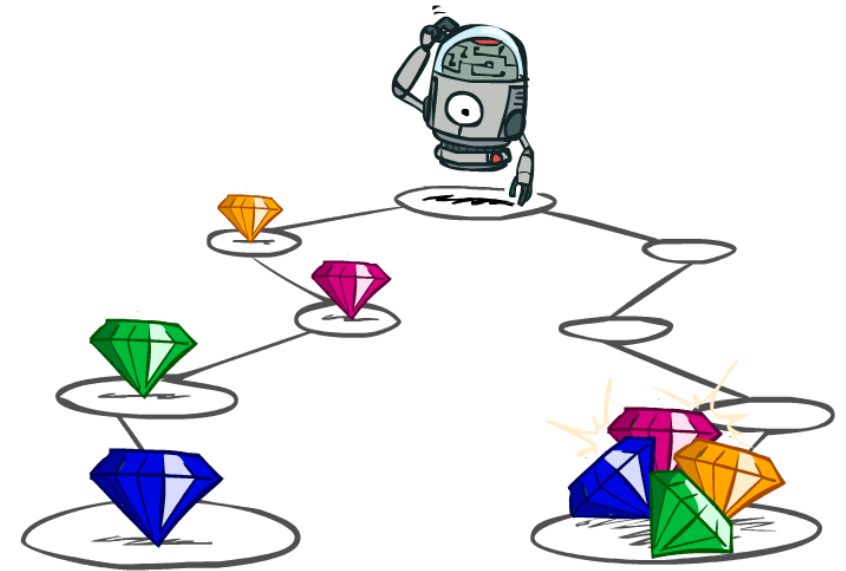
$$-0.0221 < R(s) < 0$$



$$R(s) > 0$$

Utility of the agent

- **Utility of the agent**: depends on the sequence of states, i.e., environment history $[s_0, s_1, \dots, s_n]$, rather than on a single state
 - Utility function: $U_h([s_0, s_1, \dots, s_n])$
- How to calculate utilities for state sequences (in order to be able to select the best)?
 - More or less? $[1, 2, 2]$ or $[2, 3, 4]$
 - Now or later? $[0, 0, 1]$ or $[1, 0, 0]$
- It is reasonable to maximize the sum of rewards
- It is also reasonable to prefer rewards now to rewards later
- Idea → **discount factor** $\gamma \in [0, 1]$



Discounting

- Utility is a sum of discounted rewards:

$$U[s_0, s_1, \dots] = R(s_0) + \gamma R(s_1) + \gamma^2 R(s_2) + \dots + \gamma^n R(s_n)$$

- The order of rewards matters


- Example: discount of 0.5

- $U([1,2,3]) = 1 + 0.5*2 + 0.25*3$
- $U([3,2,1]) = 3 + 0.5*2 + 0.25*1$
- $U([1,2,3]) < U([3,2,1])$

- Smaller γ means shorter-term focus

- Example: discount of 0.1


- $U([1,2,3]) = 1 + 0.1*2 + 0.01*3$



Why discounting

- Sooner rewards probably do have higher utility than later rewards
- Helps with the infinite sequences problem
 - What if the game lasts forever (infinite sequence)? Do we get infinite utilities?
 - It is also hard to compare state sequences with infinite utilities
 - It can be shown that if $\gamma < 1$ and rewards are bounded by $+/-R_{\max}$

$$U([r_0, \dots, r_\infty]) = \sum_{t=0}^{\infty} \gamma^t r_t \leq \sum_{t=0}^{\infty} \gamma^t R_{\max} = \frac{R_{\max}}{1 - \gamma}$$



- Also helps our algorithms to converge

Outline

- Introduction
- MDP formulation
- Solving the MDP
- From value-iteration to policy-iteration
- Relaxing the (full) MDP assumptions → RL
- Model-based learning
- Things you should know from this lecture & reading material

Solving MDPs

- **Input:** the MDP formulation

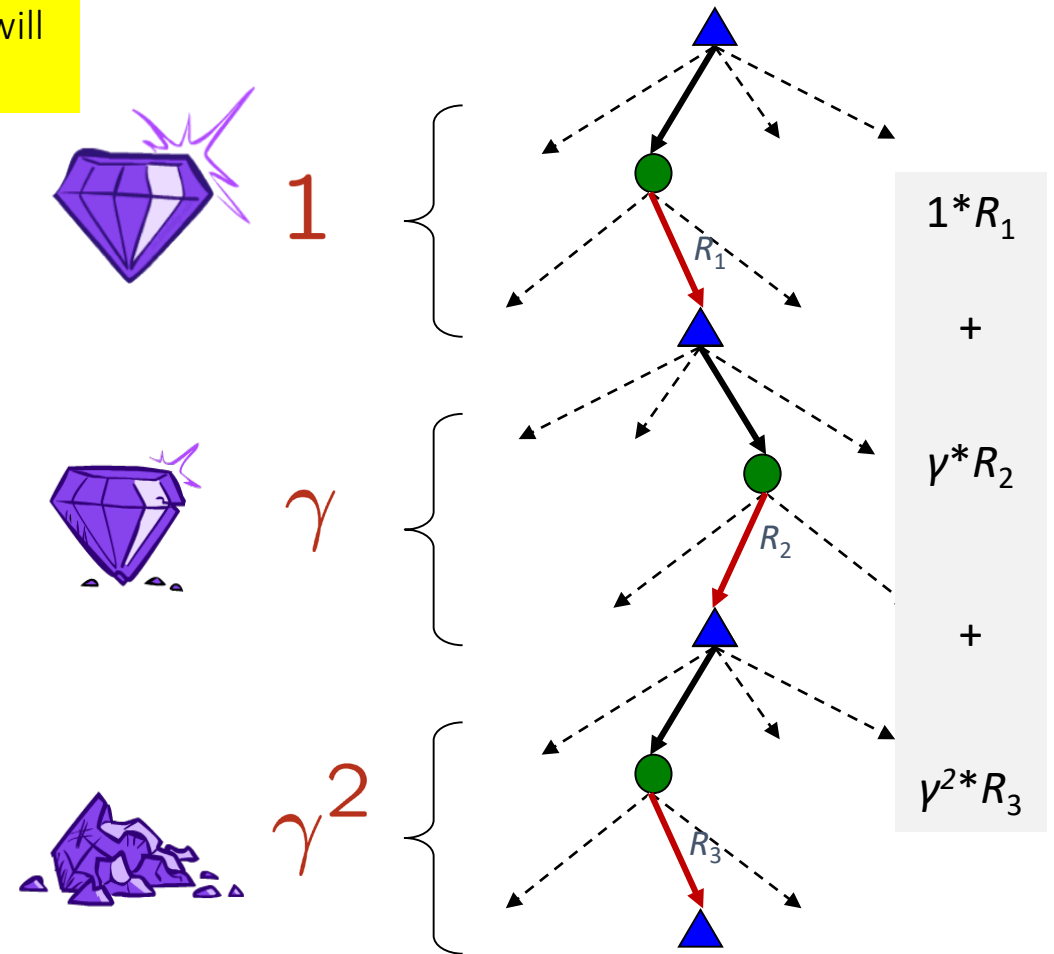
- Set of states S
- Start state S_0
- Set of actions A
- Transitions $P(s'/s, a)$ (or $T(s, a, s')$)
- Rewards $R(s, a, s')$ (and discount γ)

For the moment, we assume a full model of the world - We will relax this definition later

- **Output**

- An optimal policy: $\pi^*: S \rightarrow A$ that maps each state to an action, $\pi^*(s)$
- If followed by the agent π^* will yield the maximum expected total reward or utility

- **Utility** = sum of (discounted rewards)




Value function (A major component of an RL agent)

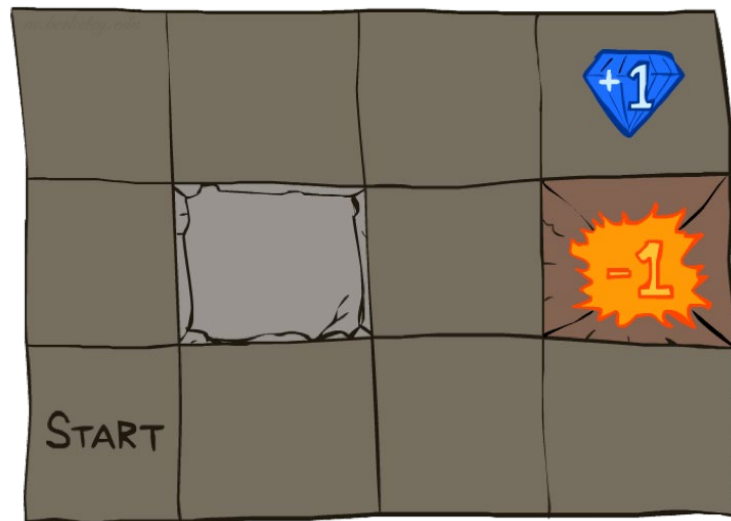
- How we decide among possible actions/states?
- The value(utility) of a state $s \rightarrow V(s)$ (called **V-value**)
 - It is a prediction of future reward
 - Used to evaluate the goodness/badness of states

- It is the expected value of the state

$$v_{\pi}(s) = \mathbb{E}_{\pi} [R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots \mid S_t = s]$$





Gridworld example: V values - utilities of states s

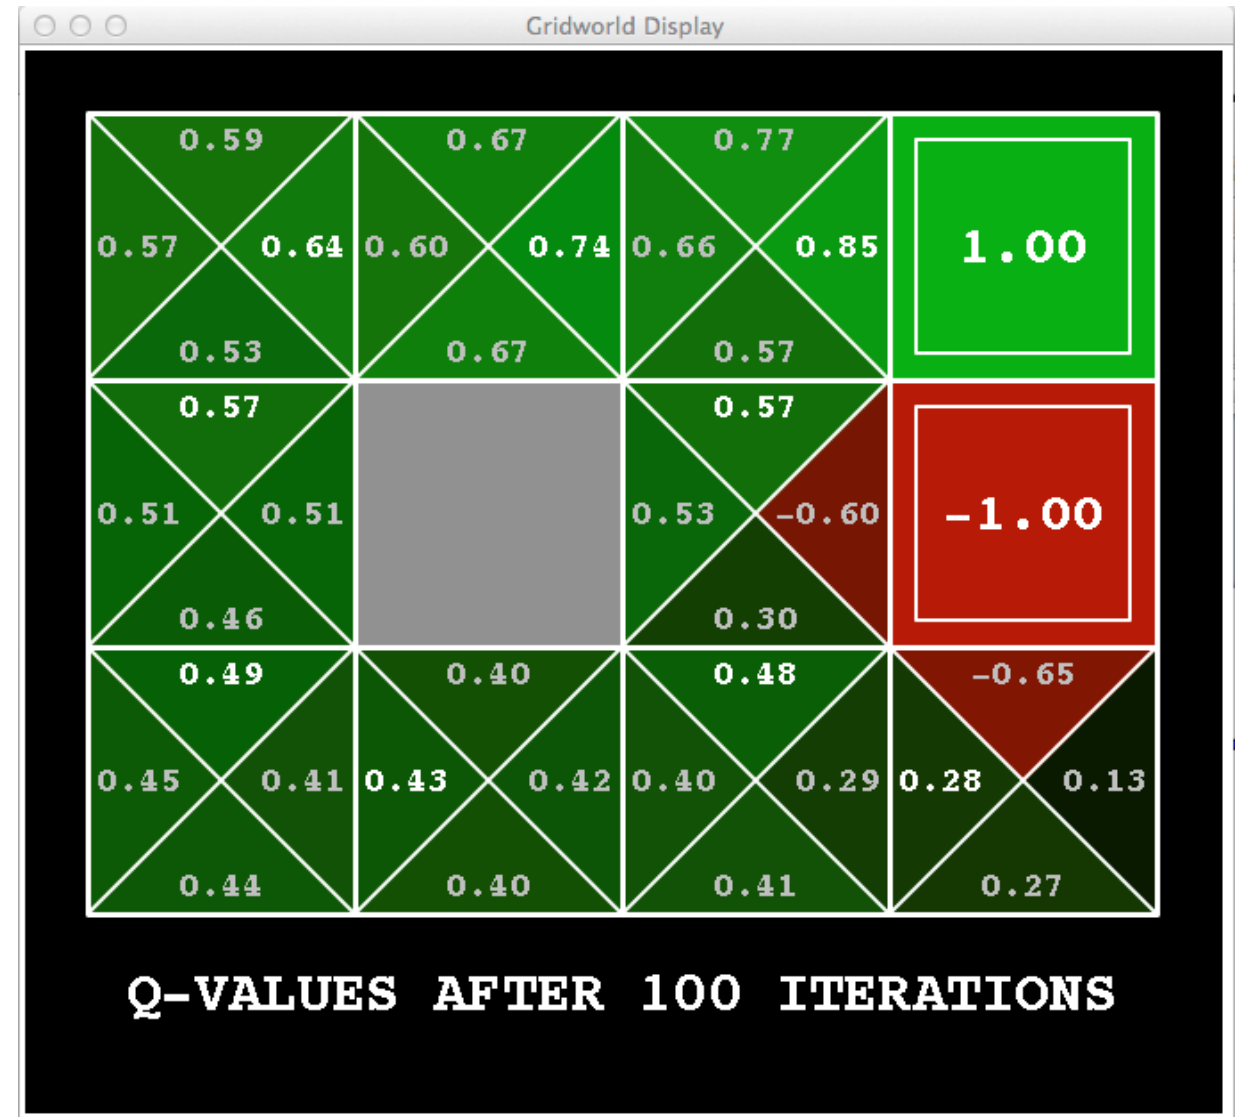


Q-states

- The combination of a state s and action a , denoted by (s,a) , is called a **Q-state**
 - It represents being in state s and having taken action a
 - still, due to uncertainty, we don't know what the outcome of the action will be
- The value of a q-state is called **Q-value**




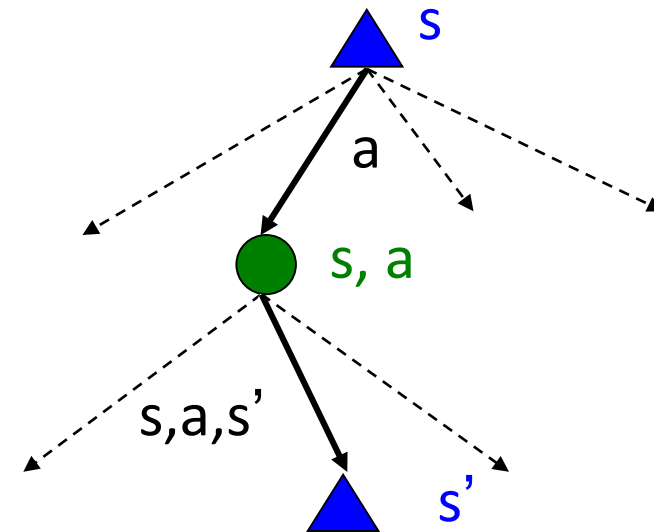
Gridworld example: Q values - utilities of states (s,a)



Optimal V^* and Q^* quantities

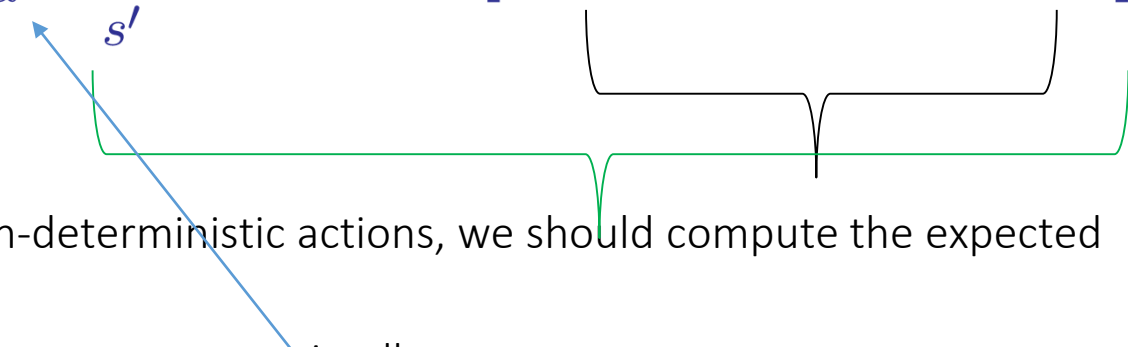
- The value (utility) of a state s :
 - $V^*(s)$ = expected utility starting in s and acting optimally from that point onwards
- The value (utility) of a q-state (s,a) :
 - $Q^*(s,a)$ = expected utility starting out having taken action a from state s and acting optimally from that point onwards
- Acting optimally: a rational agent should choose the action that maximizes the expected utility of the subsequent state
 - $\pi^*(s)$ = optimal action from state s

$$\pi^*(s) = \operatorname{argmax}_a Q^*(s, a)$$



Bellman equations

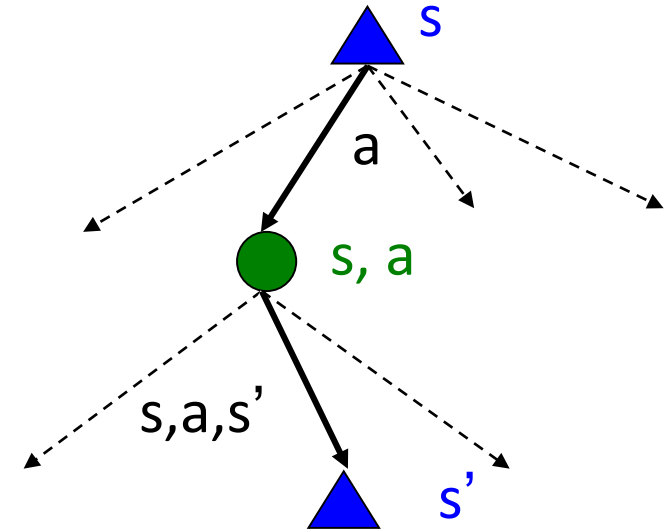
- The utility of a state is the **immediate reward** $R(s,a,s')$ for that state plus the **expected discounted utility of the next state** $\gamma V^*(s')$, assuming that the agent chooses the optimal action.

$$V^*(s) = \max_a \sum_{s'} T(s, a, s') \left[R(s, a, s') + \gamma V^*(s') \right]$$


- Due to the non-deterministic actions, we should compute the expected value
- We assume the agent acts optimally


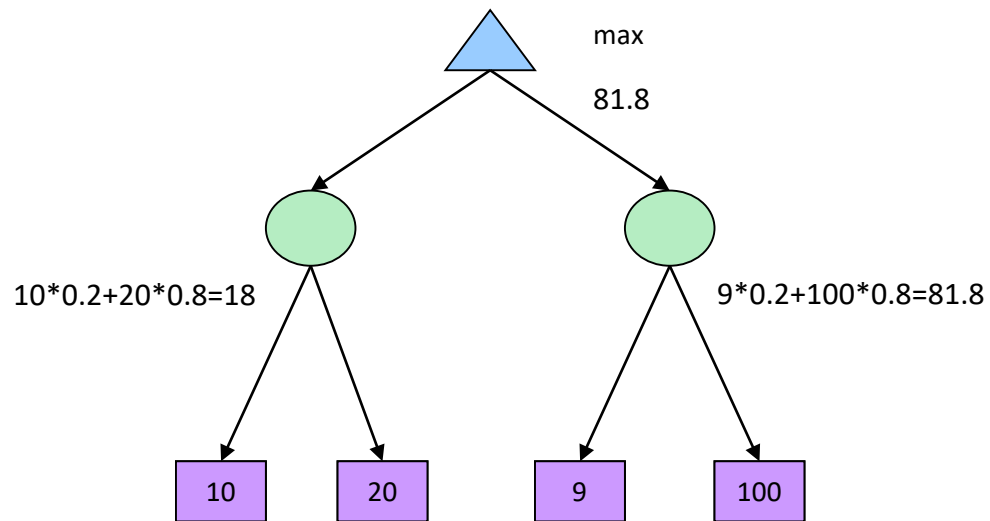
$$Q^*(s, a) = \sum_{s'} T(s, a, s') \left[R(s, a, s') + \gamma V^*(s') \right]$$

$$V^*(s) = \max_a Q^*(s, a)$$



V- and Q- values

- Small example assuming $\gamma=1, R=0$



$$V^*(s) = \max_a Q^*(s, a)$$

$$Q^*(s, a) = \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$$

The value Iteration algorithm

- The Bellman equation is the basis of the **value iteration algorithm** for solving MDPs
- If there are n possible states, then there are n Bellman equations, one for each state.

$$V^*(s) = \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$$




- We want to solve them simultaneously, but they are non-linear (max operation)
- Try an **iterative approach**

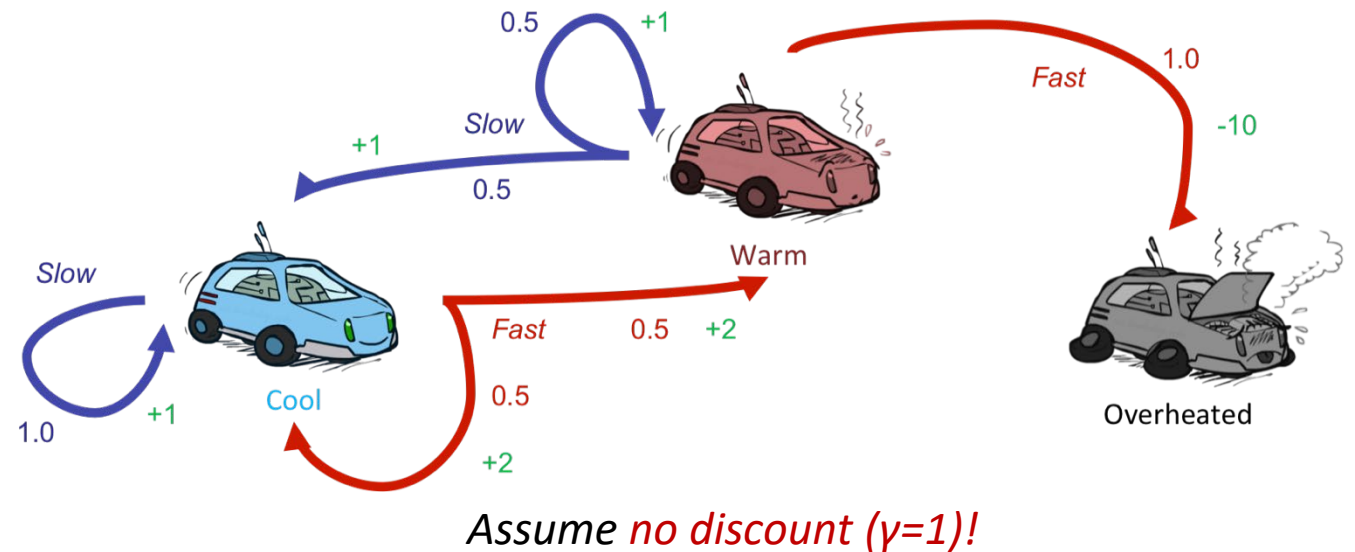
- start with arbitrary initial values for the utilities
- At each iteration k+1, for all states s
 - Update $V_{k+1}(s)$ based on $V_k(s)$ → **Bellman update**

$$V_{k+1}(s) \leftarrow \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V_k(s')]$$

- Repeat this until convergence

Example: Value Iteration

			
V_2	3.5	2.5	0
V_1	2	1	0
	0	0	0

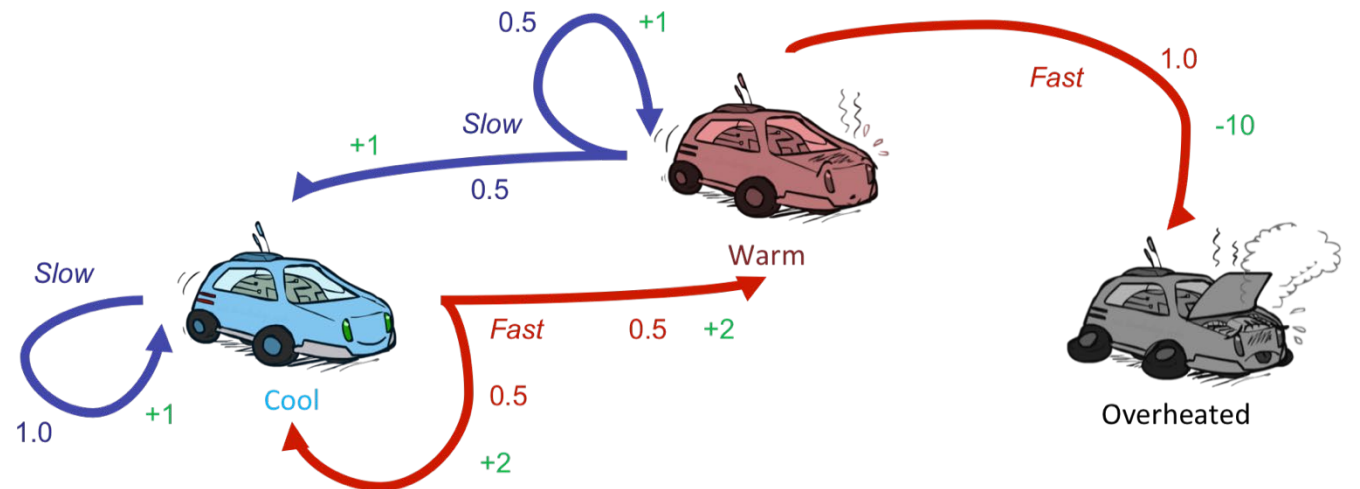
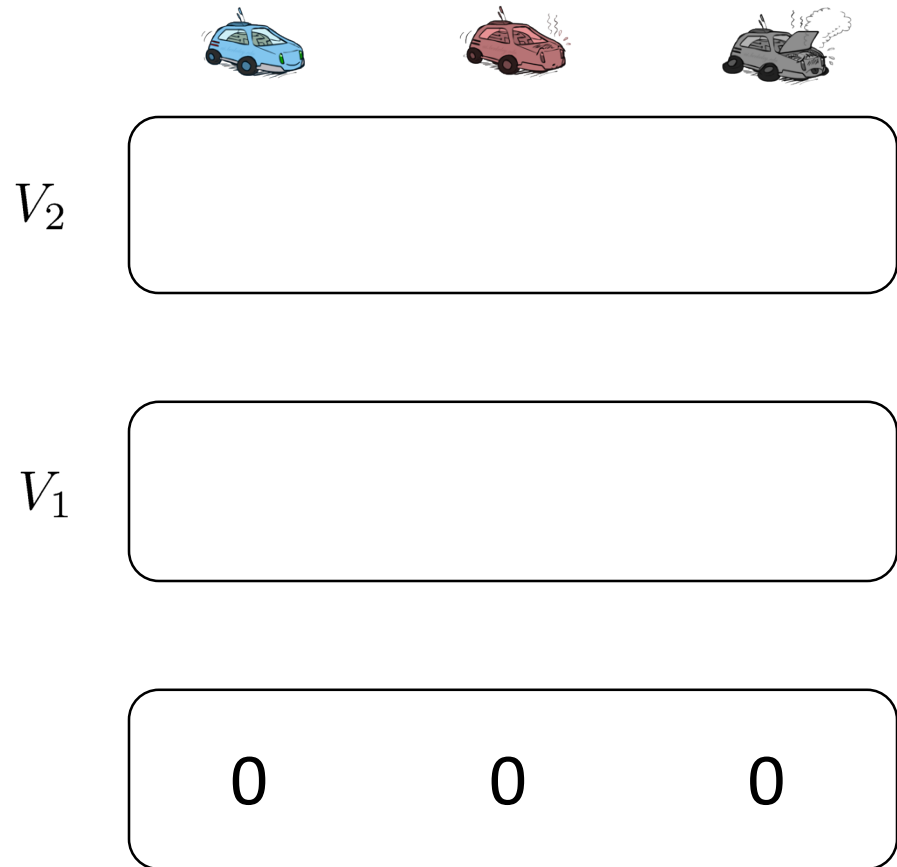


$$V_{k+1}(s) \leftarrow \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V_k(s')]$$

$$V_1(\text{cool}) = \max\{1 * [1+0], 0.5 * [2+0] + 0.5 * [2+0]\} = 2$$

$$V_2(\text{cool}) = \max\{1 * [1+2], 0.5 * [2+2] + 0.5 * [2+1]\} = \max\{3, 3, 5\} = 3.5$$

Example: Value Iteration

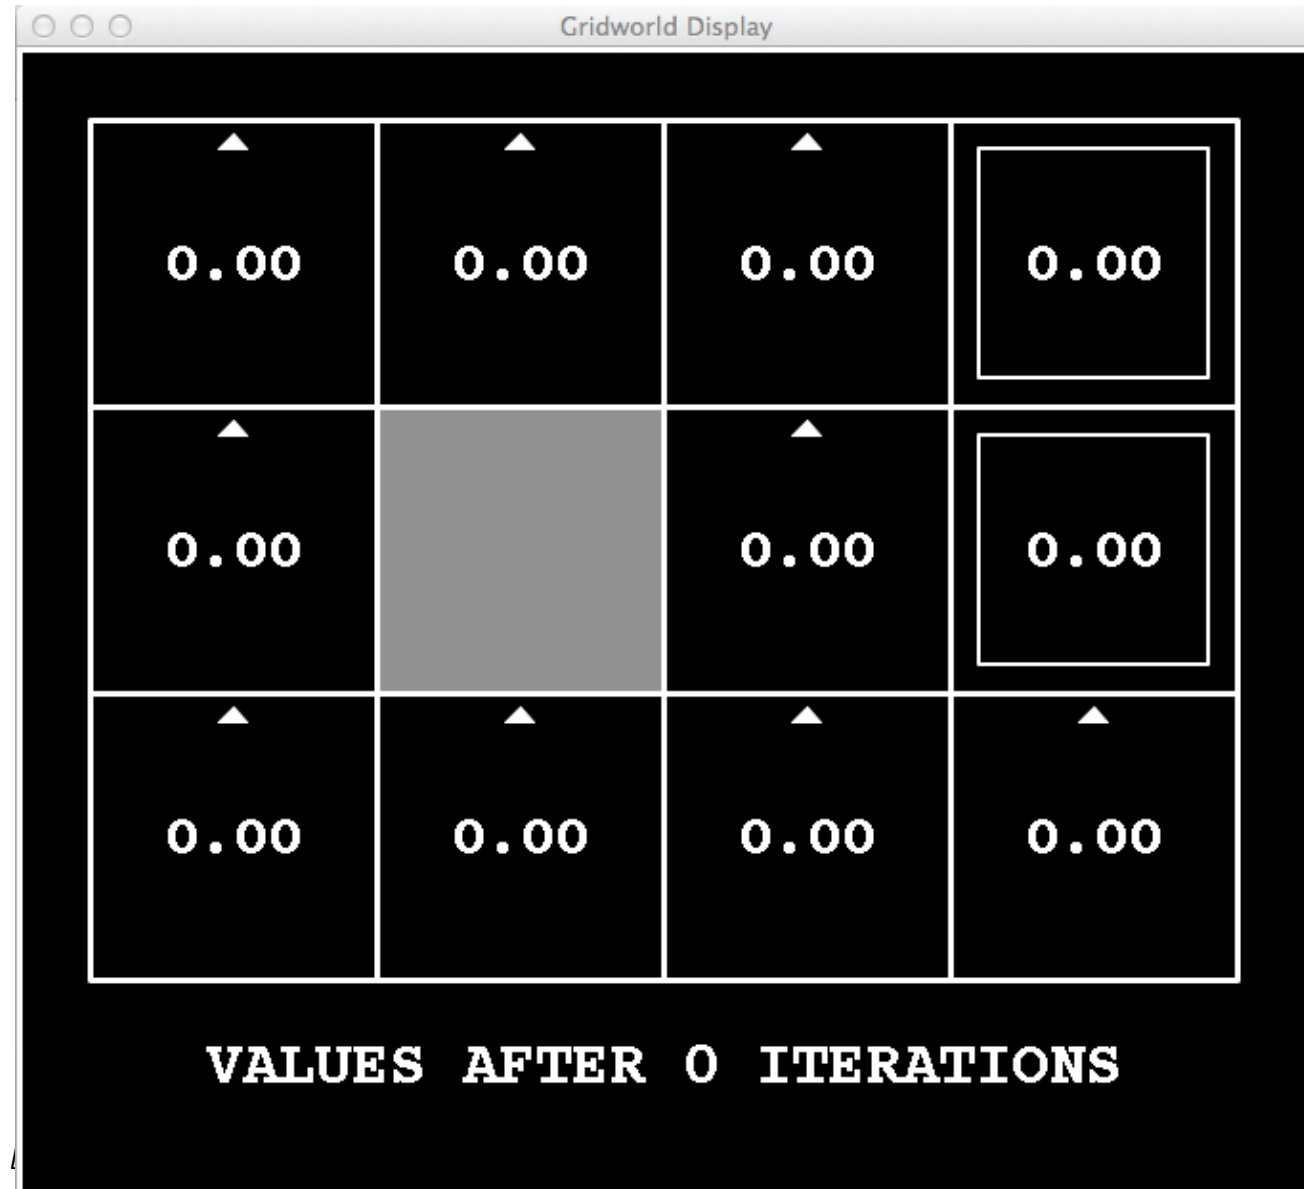


Assume a *discount* $\gamma=0.5$!

$$V_{k+1}(s) \leftarrow \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V_k(s')]$$

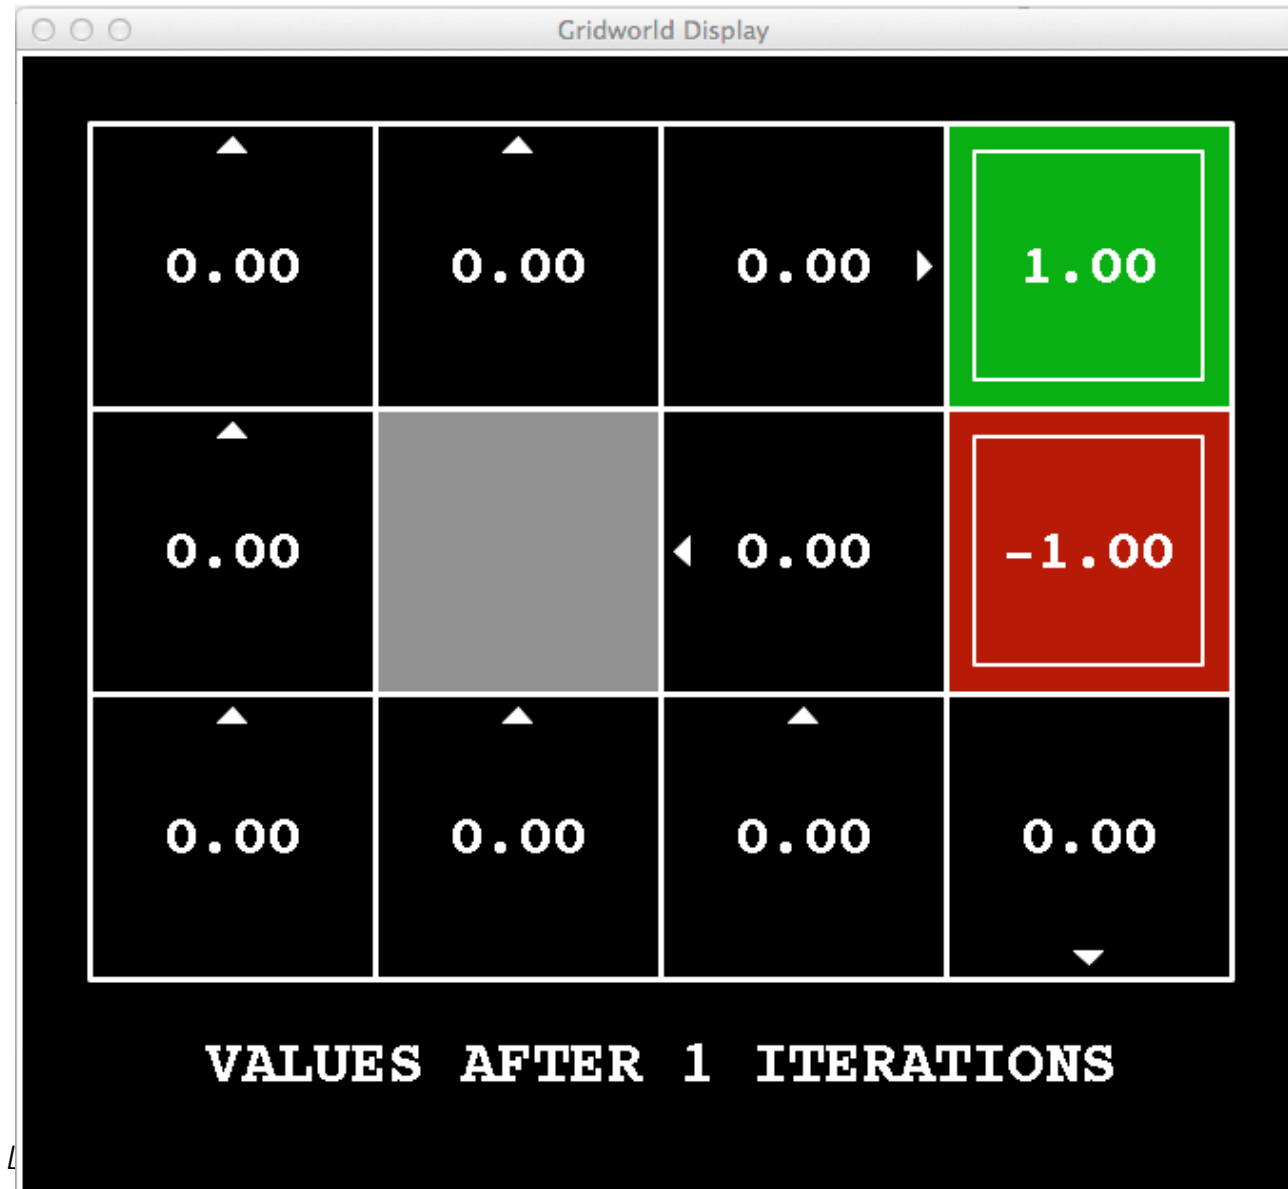
Homework!

Example: $k=0$



Noise = 0.2
Discount = 0.9
Living reward = 0

Example: $k=1$



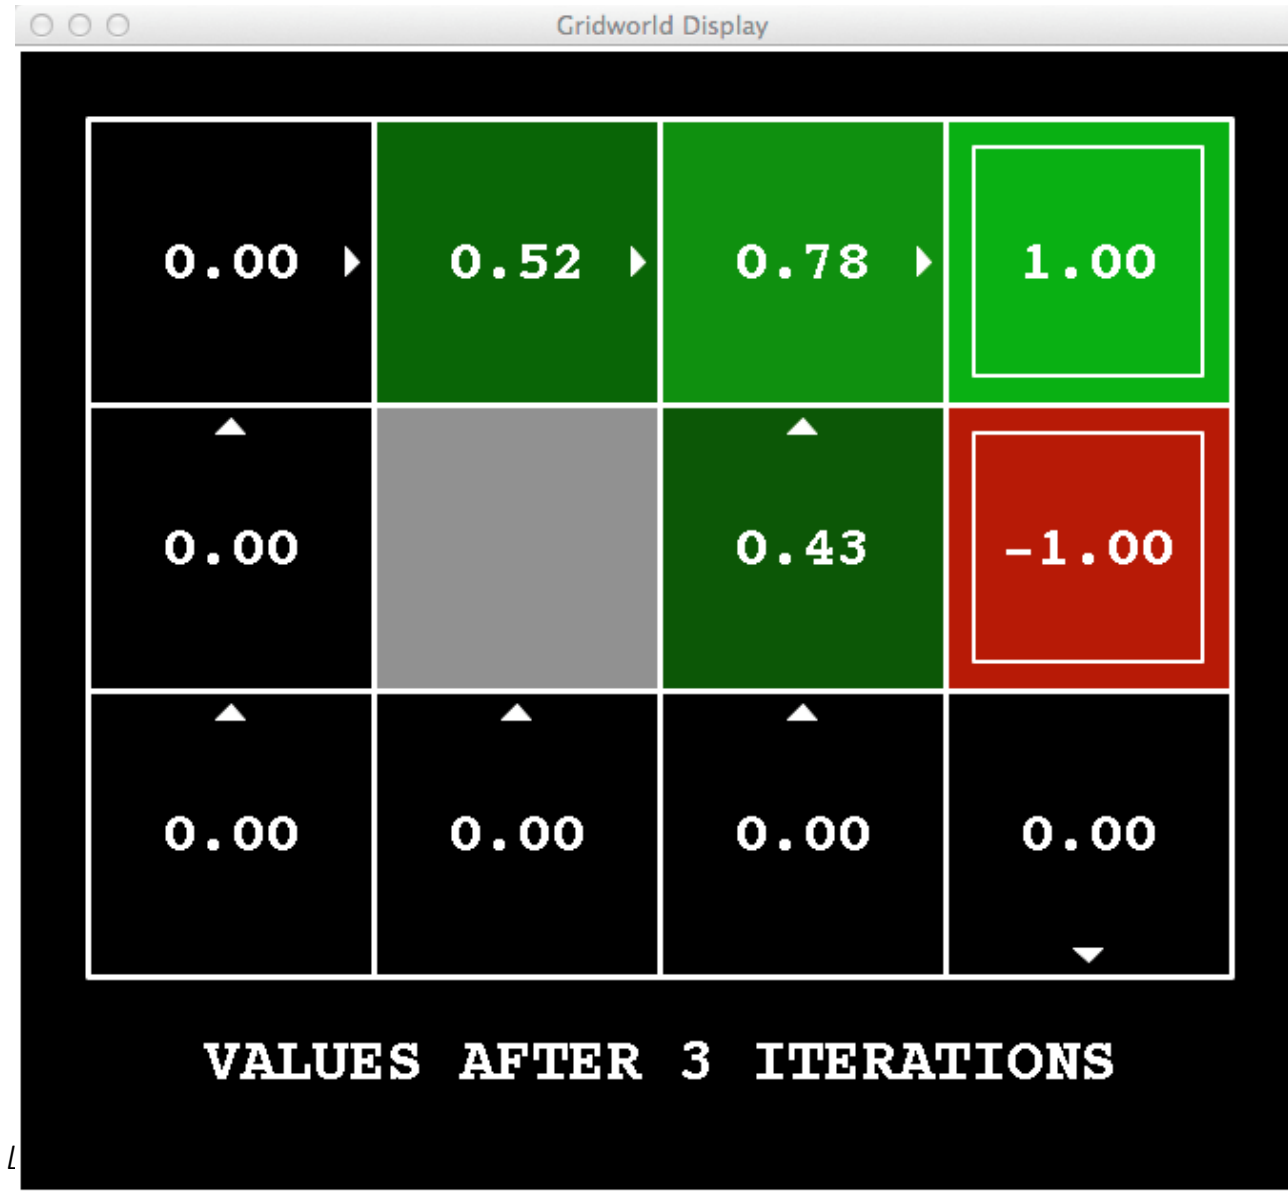
Noise = 0.2
Discount = 0.9
Living reward = 0

Example: $k=2$



Noise = 0.2
Discount = 0.9
Living reward = 0

Example: $k=3$



Noise = 0.2
Discount = 0.9
Living reward = 0

Example: $k=4$



Noise = 0.2
Discount = 0.9
Living reward = 0

Example: $k=5$



Noise = 0.2
Discount = 0.9
Living reward = 0

Example: $k=6$



Noise = 0.2
Discount = 0.9
Living reward = 0

Example: $k=7$



Noise = 0.2
Discount = 0.9
Living reward = 0

Example: $k=8$



Noise = 0.2
Discount = 0.9
Living reward = 0

Example: $k=9$



Noise = 0.2
Discount = 0.9
Living reward = 0

Example: $k=10$




Noise = 0.2
Discount = 0.9
Living reward = 0

Example: $k=11$



Noise = 0.2
Discount = 0.9
Living reward = 0

Example: $k=12$



Noise = 0.2
Discount = 0.9
Living reward = 0

Example: $k=100$

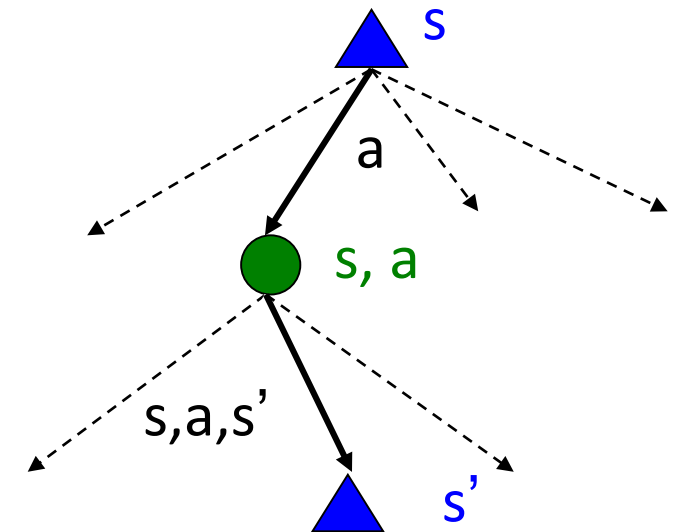


Noise = 0.2
Discount = 0.9
Living reward = 0

From optimal V^* values to optimal policy $\pi^*(s)$

- At convergence we find $V^*(s)$ based on which we can find policy $\pi^*(s)$
- **Policy extraction**: given the optimal values, what is the implied optimal policy?
- The V -values are non-actionable, we need to look 1-step ahead

$$\begin{aligned}\pi^*(s) &= \arg \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')] \\ &= \arg \max_a Q^*(s, a)\end{aligned}$$

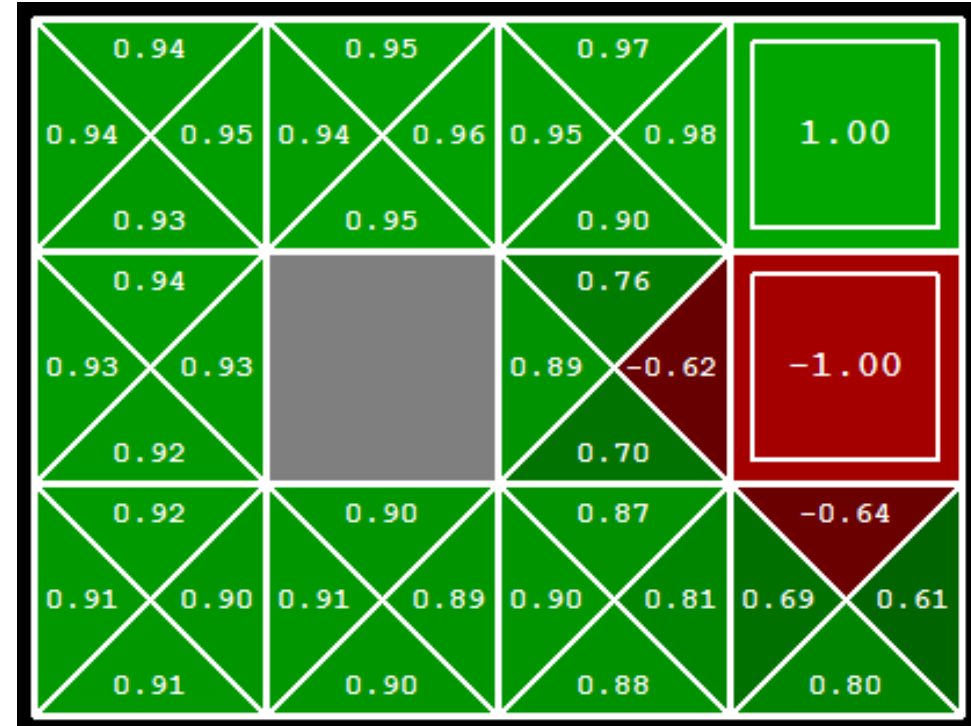


Policy extraction from Q^* -values

- If we have q -values, completely trivial to decide
 - Select the action that takes us to the q -state with the max q -value

$$\pi^*(s) = \arg \max_a Q^*(s, a)$$

- So, it is better to keep the q -values



Outline

- Introduction
- MDP formulation
- Solving the MDP
- From value-iteration to policy-iteration
- Relaxing the (full) MDP assumptions → RL
- Model-based learning
- Things you should know from this lecture & reading material

From value-iteration to policy-iteration

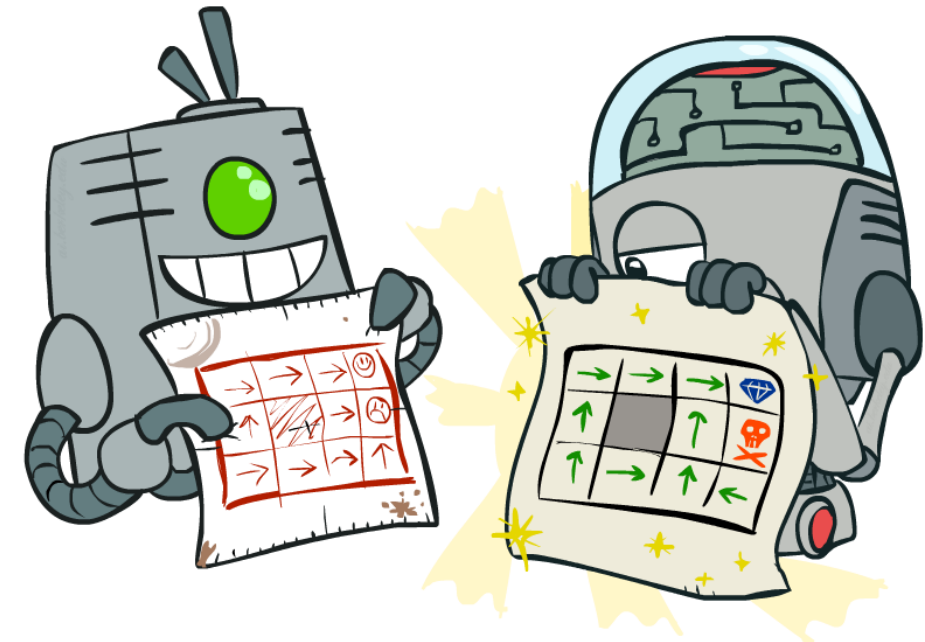
- In value iteration, approximations get refined towards optimal values
- But value convergence takes too long
- Policy might converge faster than values



- So, it is possible to get an optimal policy even if the utility function estimates are inaccurate



Policy iteration

- The policy iteration algorithm consists of two steps
 - **Step 1. Policy evaluation:** calculate utilities for some fixed policy $\pi()$ (not optimal utilities!) until convergence
 - **Step 2. Policy improvement:** update policy using one-step look-ahead with resulting converged (but not optimal!) utilities as future values
- Repeat steps 1, 2 until convergence



Step 1: Policy evaluation

- Given a **fixed policy** π calculate the utility of each state s if π were to be executed
 - $V^\pi(s)$: the utility of s according to π

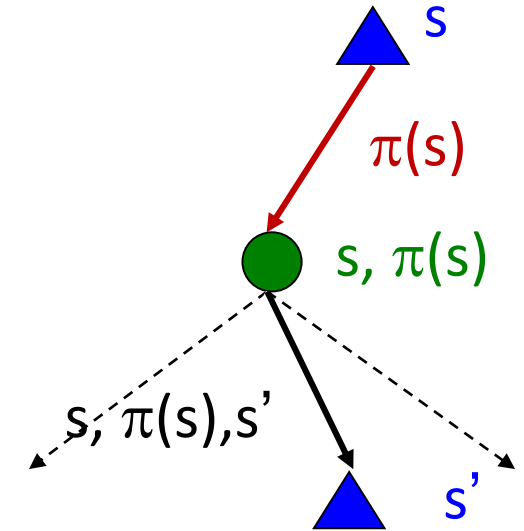


- For the optimal action we need to take **max** over all actions to compute the optimal values
- If we have a fixed policy $\pi()$, we only need **one** action per state $\pi(s)$
 - of course, the result depends on which policy we fixed

Step 1: Policy evaluation

- Define the utility of a state s under a fixed policy π .
 - $V_{\pi}(s)$ = expected total discounted rewards starting in s and following π
- Recursive relation (one-step look-ahead / **simplified Bellman equation**):

$$V^{\pi}(s) = \sum_{s'} T(s, \pi(s), s') [R(s, \pi(s), s') + \gamma V^{\pi}(s')]$$



compare it to optimal policy approach

$$V^*(s) = \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$$

Step 1: Policy evaluation - computation

- How do we calculate the V 's for a fixed policy π ?


$$V^\pi(s) = \sum_{s'} T(s, \pi(s), s') [R(s, \pi(s), s') + \gamma V^\pi(s')]$$

- Idea 1: Turn recursive Bellman equations into updates (like value iteration algorithm)

$$V_0^\pi(s) = 0$$

$$V_{k+1}^\pi(s) \leftarrow \sum_{s'} T(s, \pi(s), s') [R(s, \pi(s), s') + \gamma V_k^\pi(s')]$$

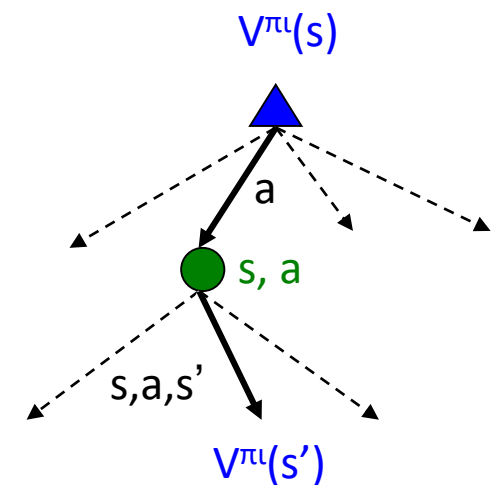
- Much faster than the value iteration approach
- Idea 2: Without the max operation, the Bellman equations are just a linear system.
 - We can solve them using exact solution methods (if state space is small)



Step 2: Policy Improvement

- We can evaluate a fixed policy π (using **policy evaluation**) $\rightarrow V^\pi(s)$
- How can we improve π ?
- **Policy improvement**: with fixed values $V_\pi(s)$, find the best action according to one-step-look-ahead (so, use policy extraction)

$$\pi_{i+1}(s) = \arg \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^{\pi_i}(s')]$$



Policy Iteration algorithm

- **Step 1 – Policy evaluation:** with fixed current policy π , find values with policy evaluation

- Iterate until values converge (simplified Bellman update formula):

$$V_{k+1}^{\pi_i}(s) \leftarrow \sum_{s'} T(s, \pi_i(s), s') [R(s, \pi_i(s), s') + \gamma V_k^{\pi_i}(s')]$$

- Note: could also solve value equations with other techniques

- **Step 2 – Policy improvement:** with fixed values, find the best action according to one-step-look-head (so, use policy extraction)

$$\pi_{i+1}(s) = \arg \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^{\pi_i}(s')]$$

- **Repeat steps 1, 2 until convergence**

- This is policy iteration

- It's still optimal!
- Can converge (much) faster (under some conditions)

Summary: MDP Algorithms

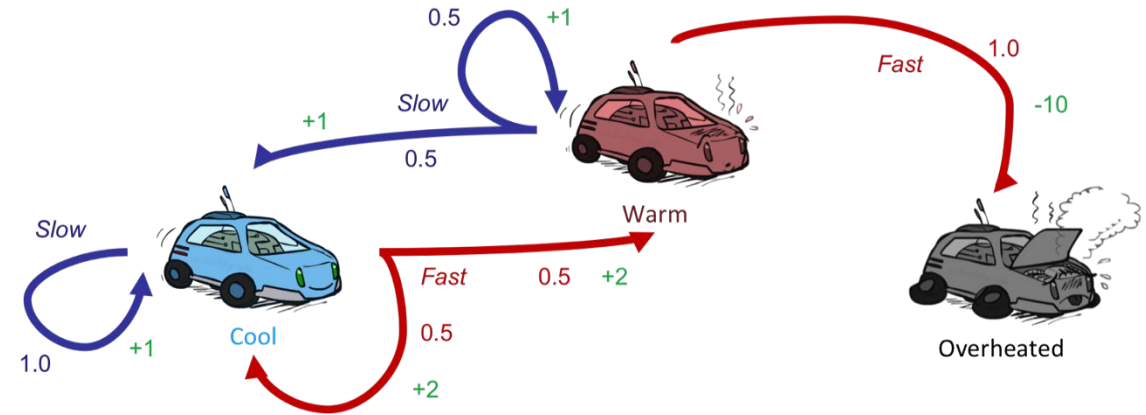
- So you want to....
 - Compute optimal values: use **value iteration** or **policy iteration**
 - Compute values for a particular policy: use **policy evaluation**
 - Turn your values into a policy: use **policy extraction**
- These all look the same!
 - They basically are – they are all variations of Bellman updates
 - They differ only in whether we plug in a fixed policy or max over actions

Outline

- Introduction
- MDP formulation
- Solving the MDP
- From value-iteration to policy-iteration
- Relaxing the (full) MDP assumptions → RL
- Model-based learning
- Things you should know from this lecture & reading material

From MDPs to Reinforcement Learning

- In a MDP, we have
 - A set of states $s \in S$
 - A set of actions (per state) A
 - A transition model $T(s,a,s')$
 - A reward function $R(s,a,s')$
- and we are looking for a policy $\pi(s)$



From MDPs to Reinforcement Learning

- In Reinforcement Learning (RL)

- We still have an MDP

- A set of states $s \in S$
 - A set of actions (per state) A
 - A transition model $T(s,a,s')$
 - A reward function $R(s,a,s')$

- Still looking for a policy $\pi(s)$



- New twist: we don't know T , R

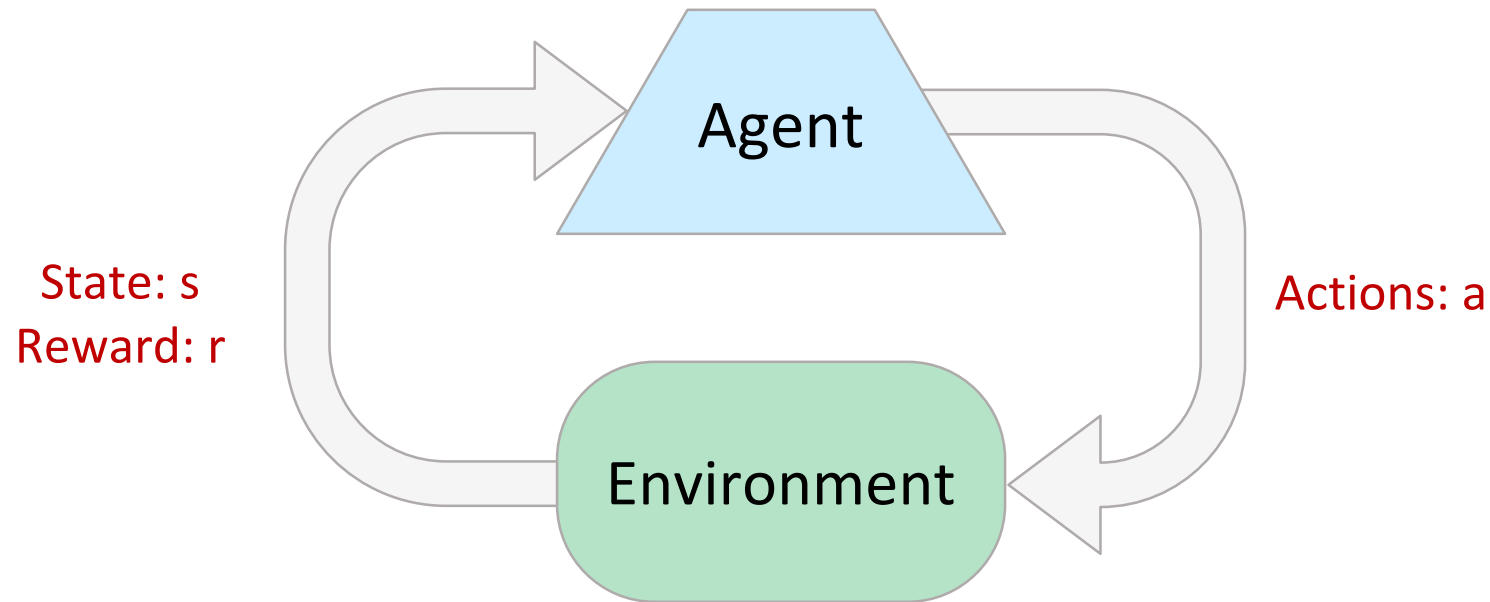
- i.e. we don't know which states are good or what the actions do

- So, we must actually try out actions and states to learn

So RL can solve MDP problems
when we don't know the MDP

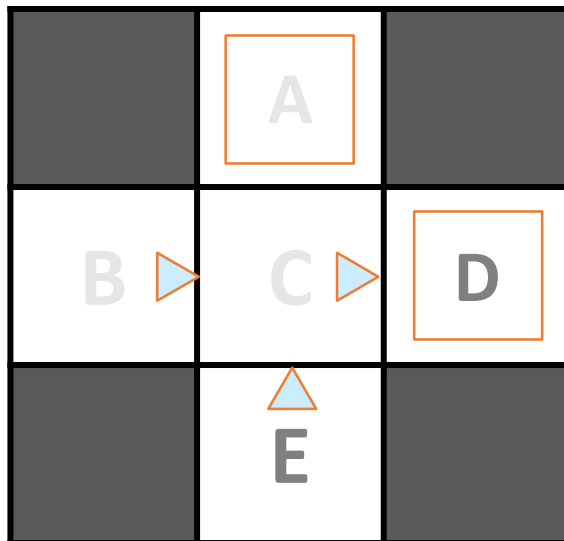
Reinforcement Learning

- Basic idea:
 - Agent receives feedback in the form of rewards
 - Agent's utility is defined by the reward function
 - Must (learn to) act so as to maximize expected utility
 - All learning is based on observed samples of outcomes!



Collecting experience

- The agent collects **experience/data** via its interaction with the environment
- Tuples (s, a, s', r) are known as **samples**
- A collection of samples until arriving at a terminal state is known as **episode**



Observed Episodes (Training)

Episode 1

B, east, C, -1
C, east, D, -1
D, exit, x, +10

Episode 2

B, east, C, -1
C, east, D, -1
D, exit, x, +10

Episode 3

E, north, C, -1
C, east, D, -1
D, exit, x, +10

Episode 4

E, north, C, -1
C, east, A, -1
A, exit, x, -10

Key ideas for learning

- Online vs Offline/Batch learning

- Learn while exploring the world, or learn from fixed batch of data

- Active vs. Passive Learning

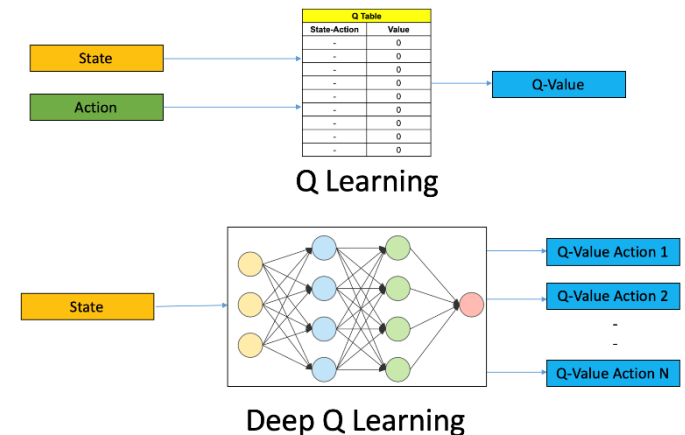
- Does the learner actively choose actions to gather experience? or, is a fixed policy provided?

- Model-based vs. Model-free Learning

- Do we estimate $T(s,a,s')$ and $R(s,a,s')$, or just learn values/policy directly?

- What we will (quickly ☹) cover in the next 3-4 lectures

- Model-based learning
 - Model-free learning
 - Passive RL (direct evaluation, TD-learning)
 - Active RL (Q-learning)
 - Value-function approximation (Approximate (deep) Q-learning)



Outline

- Introduction
- MDP formulation
- Solving the MDP
- From value-iteration to policy-iteration
- Relaxing the (full) MDP assumptions → RL
- Model-based learning
- Things you should know from this lecture & reading material

Model-based learning

- Model-based idea:
 - (Step 1) Learn an **approximate** model of T, R based on experiences/data
 - (Step 2) Solve the MDP based on the **learned T, R**
- **Step 1:** Learn **empirical** MDP model
 - Count outcomes s' for each q-state (s, a)
 - Normalize to give an estimate of $\hat{T}(s, a, s')$
 - Discover each $\hat{R}(s, a, s')$ estimate when we experience (s, a, s')
- **Step 2:** Solve the learned MDP
 - For example, use value iteration or policy iteration (see previous slides)


Example: Model-based learning – Step 1: Learn the empirical model

Recall that $T(s,a,s')=P(s' | s,a)$

$T(C, \text{east}, D)=P(D|C, \text{east})=3/4$

$T(C, \text{east}, A)=P(D|C, \text{east})=1/4$

Input Policy π



Assume: $\gamma = 1$

Observed Episodes (Training)

Episode 1

B, east, C, -1
C, east, D, -1
D, exit, x, +10

Episode 2

B, east, C, -1
C, east, D, -1
D, exit, x, +10

Episode 3

E, north, C, -1
C, east, D, -1
D, exit, x, +10

Episode 4

E, north, C, -1
C, east, A, -1
A, exit, x, -10

Step 1:

For each q-state, count outcomes, e.g., for (C,east):
{C, east, D, -1;
C, east, D, -1;
C, east, D, -1;
C, east, A, -1;}

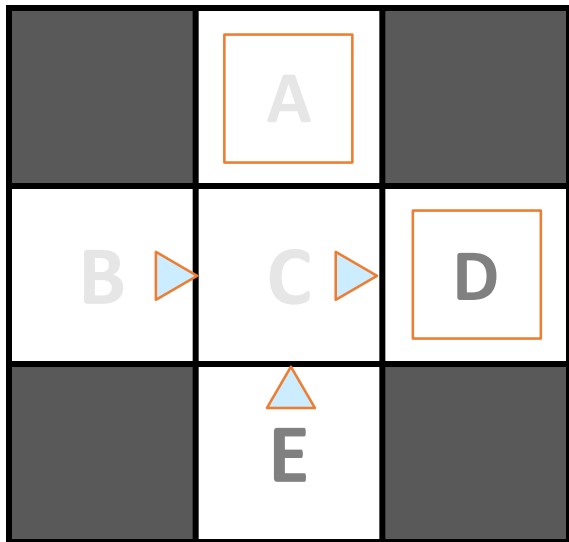
Based on these outcomes, we can compute the probability of each outcome $\rightarrow T$

For each sample (s,a,s') discover the associated reward $\rightarrow R$

Assumption: the reward is deterministic

Example: Model-based learning – Step 1: Learn the empirical model

Input Policy π



Assume: $\gamma = 1$

Observed Episodes (Training)

Episode 1

B, east, C, -1
C, east, D, -1
D, exit, x, +10

Episode 2

B, east, C, -1
C, east, D, -1
D, exit, x, +10

Episode 3

E, north, C, -1
C, east, D, -1
D, exit, x, +10

Episode 4

E, north, C, -1
C, east, A, -1
A, exit, x, -10

Learned Model

$$\hat{T}(s, a, s')$$

$T(B, \text{east}, C) = 1.00$
 $T(C, \text{east}, D) = 0.75$
 $T(C, \text{east}, A) = 0.25$
...

$$\hat{R}(s, a, s')$$

$R(B, \text{east}, C) = -1$
 $R(C, \text{east}, D) = -1$
 $R(D, \text{exit}, x) = +10$
...

Assumption: the reward is deterministic

Example: Model-based learning – Step 2: Solve the learned MDP

- As the experience increases and we collect more and more samples, our empirical models will improve:
 - The estimated transition model \hat{T} will converge towards real T
 - new rewards will be discovered as new (s,a,s') tuples are explored.
- When the estimates are adequate, the training phase (Step 1) ends
- (Step 2)Based on the learned parameters, we solve the (conventional) MDP
 - using value iteration or policy iteration (see previous slides)

Model-based learning: discussion

- Model-based Idea:
 - (Step 1) Learn an approximate model of T, R based on experiences/data
 - (Step 2) Solve the conventional MDP based on the learned T, R
- Pros
 - Very simple and intuitive
 - Remarkably effective
- Cons
 - Sufficient (training) experience is required
 - Maintaining all these counts is expensive

Outline

- Introduction
 - MDP formulation
 - Solving the MDP
 - From value-iteration to policy-iteration
 - Relaxing the (full) MDP assumptions → RL
 - Model-based learning
- Things you should know from this lecture & reading material

Overview and Reading

■ Overview

- ❑ RL basics
- ❑ MDP formulation
- ❑ Bellman equations
- ❑ Value iteration
- ❑ Policy extraction
- ❑ Policy evaluation
- ❑ Policy iteration

■ Reading

- ❑ Chapter 16& 23, AI book, 4th edition
- ❑ RL Book, Barto and Sutton, 2nd edition
- ❑ [Introduction to Reinforcement Learning with David Silver](#) (DeepMind)
- ❑ [Stanford CS234: Reinforcement Learning](#) with Emma Brunskill

Hands on experience



- Small programming exercise
 - For a (small) grid-world example similar to our toy example implement from scratch
 - Value iteration
 - Policy extraction
 - Policy evaluation (try out different policies, e.g., go always in one direction {N,S,W,E}, act randomly etc)
 - Policy iteration
 - Assuming now that you don't know the R , T components, implement the model-based RL version
- Familiarize yourself with [OpenAI Gym](#)
 - Gym is a toolkit for developing and comparing reinforcement learning algorithms. It supports teaching agents everything from [walking](#) to playing games like [Pong](#) or [Pinball](#).
- We will release a 3rd project (CartPole balancing problem, most probably)
 - Look at this presentation of the [CartPole problem](#) (states, actions, rewards, ...)

Thank you

Questions/Feedback/Wishes?

Acknowledgements

- The slides are based on
 - CS 188 | Introduction to Artificial Intelligence, Berkeley
 - Artificial Intelligence: A modern approach (Russel and Norvig), 4th edition