

# Lecture: Machine Learning for Data Science

Winter semester 2021/22

Lectures 25 & 26: Velocity (stream clustering)


Prof. Dr. Eirini Ntoutsi

# Outline

- Data stream clustering basics
- Summarization
- Overview of stream clustering methods
- Partitioning methods
- Density-based methods
- Grid-based methods
- Data stream clustering: evaluation aspects
- Things you should know from this lecture & reading material

## (recap from previous week) Data streams

- “A data stream is a **potentially unbounded, ordered sequence of data items**, which arrive **continuously at high-speeds**”  
Springer Encyclopedia of Machine Learning, 2017

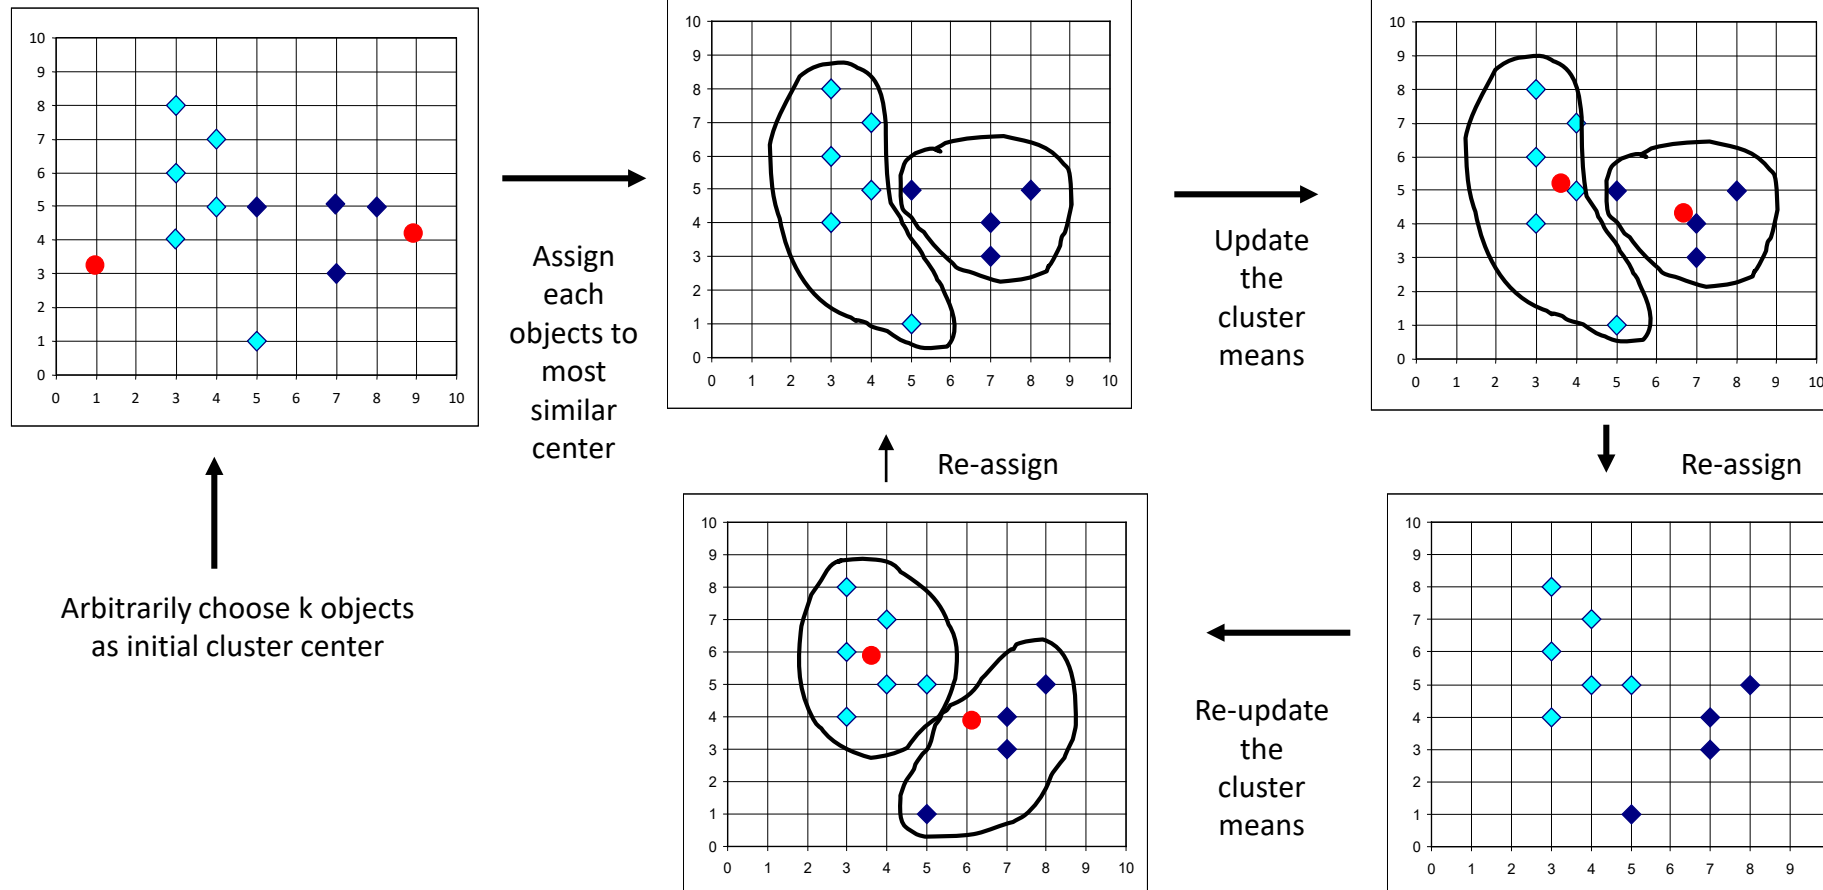


- Key characteristics
  - ❑ **Huge volumes of continuous data, possibly infinite**: Random access is expensive or undesirable (due to e.g., privacy)
  - ❑ **High arrival rate**: response time matters
  - ❑ **Non-stationary/ evolving data**: Data evolve over time as new data arrive and old data become obsolete/irrelevant

# (recap from previous week) Example: batch k-Means (see also lecture 10)

- The **complete dataset** is given as input to the algorithm
- The dataset is **accessed multiple times** during the iterations

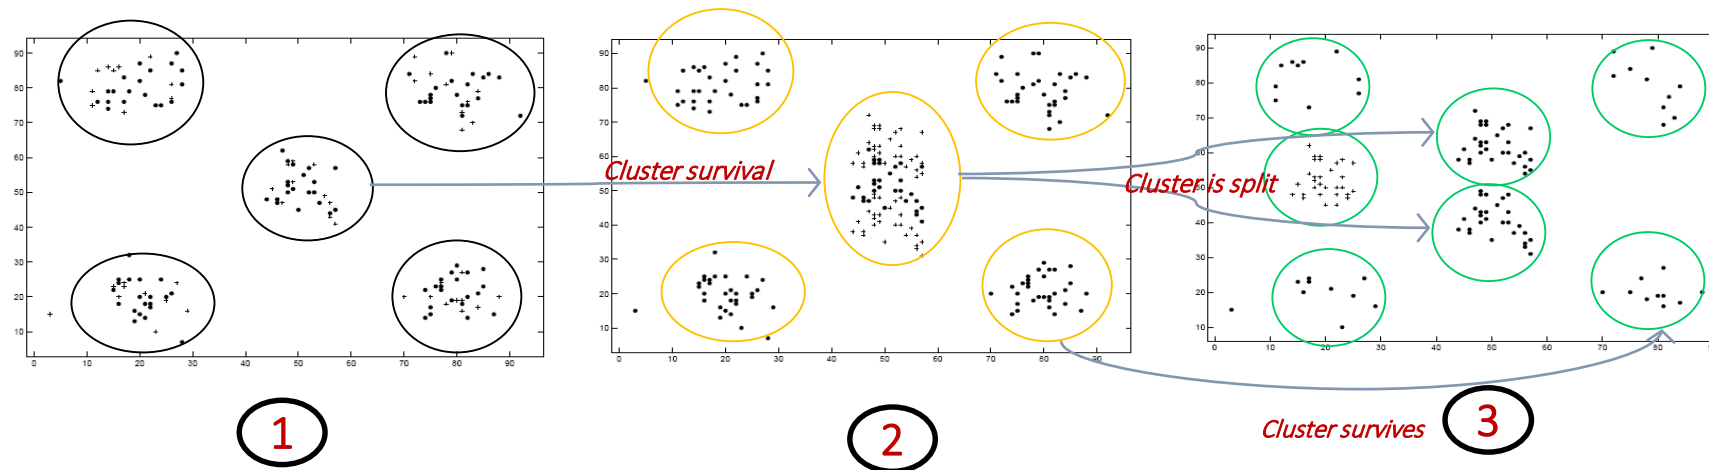
k=2



# (recap from previous week) Why the stationarity violation is problematic?

Unsupervised learning case

- As data evolve with time, the clustering is becoming invalid/obsolete
- **External changes:** the relationship of a cluster to the other clusters might change, e.g., cluster survival, split, merge, appearance, disappearance
- **Internal changes:** the description of a cluster might change both externally (i.e., cluster members) and internally (cluster properties)



Source: The MONIC framework, Spiliopoulou et al, KDD06

## (recap from previous week) Requirements for stream learning

- Need for new learning algorithms that
  - have the ability to **incorporate new data** (incremental models)
  - deal with non-stationary data generation processes
    - Ability to discard obsolete data (or, obsolete (parts of the) model) (**data ageing/ forgetting**)
- subject to:
  - **resource constraints** (processing time, memory)
  - **single scan of the data** (one look, no random access)

# Clustering data streams

- Clustering is one of the core learning tasks
  - Used as either a standalone tool or as a preprocessing tool
- The (batch) clustering problem:
  - Given a set of data instances, the goal is to group the data into groups of similar data (clusters)
  - The dataset is available from the beginning to the algorithm
  - The algorithm is allowed to iterate over the dataset
- The **data stream clustering problem**:
  - Maintain a good clustering over the (non-stationary) stream, subject to resource constraints
    - No random access to the data (you only have a look at the data instances when they arrive)
    - The processing time per instance should be low

# Challenges & Requirements for data stream clustering

- Traditional clustering methods require access upon the whole dataset
  - work with summaries, rather than raw data
- The underlying population distribution might change:
  - the clustering structure needs to be maintained/adapted online
  - one clustering model might not be adequate to capture the evolution of the underlying population
- The role of outliers and clusters are often exchanged in a stream
  - timely and accurate identification of outliers is necessary



# Outline


- Data stream clustering basics
- Summarization
- Overview of stream clustering methods
- Partitioning methods
- Density-based methods
- Grid-based methods
- Data stream clustering: evaluation aspects
- Things you should know from this lecture & reading material

# Dealing with the efficiency requirements

- In the stream classification part, we mainly focused on the non-stationary aspect and we didn't directly address efficiency!
  - For the clustering part, many of the methods directly address efficiency (and of course aim at good quality clustering results)
- **Efficiency challenge** (not restricted to clustering): How can we learn from high volumes of data faster?
- Performance depends on
  - the volume of the data set (cardinality, dimensionality)
  - the scalability of the learning algorithms
  - ...
- Solutions for speeding up learning
  - Use **high-performance computing architectures**
    - Parallel computing; Distributed computing; Cloud computing; ...
  - **Reduce the number of objects being processed**
    - **Summarization/Compression**: “compress” the data using “higher-level” descriptors (summaries) (the most relevant for our discussion on stream clustering)
    - **Sampling**: select a subset of the data to work on
    - Keep **quality data** (lately known as **data-centric AI**)
      - Andrew Ng is one of the supporters of this idea, listen e.g., [“A Chat with Andrew on MLOps: From Model-centric to Data-centric AI”](#)
  - **Develop more efficient methods** 😊
  - ....

# Summarization/Compression

- Summarization/Compression is one way to speed up learning
- Main idea:
  - Summarize/Compress the input data into a set of summaries. Original/ raw data are discarded.
  - Apply machine learning algorithms upon the summaries afterwards
- Why does **summarization** makes sense?
  - Summaries comprise lossy but still good representations of the original raw data
  - Having good summaries is of course critical!



# Summarization/Compression

- Examples of summaries
  - Cluster feature vectors [Zhang et al 1996]
  - Data bubbles [Breuning et al 2001]
  - ...
- Again, the quality of the summaries is of paramount importance for the success of the learning task

# Cluster-feature (CF) vectors/ BIRCH algorithm

- **BIRCH** (Balanced Iterative Reducing and Clustering using Hierarchies) [Zhang et al 1996]
  - BIRCH is the first approach for clustering large scale data
- BIRCH introduced the idea of **cluster feature vector summaries**
  - also known as **microclusters** in the stream clustering domain
- BIRCH organizes the CF summaries into a tree structure
  - **CF tree**: A multi-level compression of the data that tries to preserve the inherent clustering structure of the data

# Cluster feature vector (CF) summaries or micro-clusters

Given  $N$   $d$ -dimensional points in a cluster  $C$ , the cluster feature (CF) vector of  $C$  is defined as a triple:

$$CF = (N, \overrightarrow{LS}, SS)$$

where

□  $N = |C|$  is the number of points in  $C$

LS stands for Linear Sum  
SS stands for Square Sum

□  $\overrightarrow{LS} = \sum_{i=1}^N \vec{X}_i$  is the linear sum of the  $N$  data points

□  $SS = \sum_{i=1}^N \vec{X}_i^2 = \sum_{i=1}^N \langle X_i, X_i \rangle$  is the square sum of the  $N$  data points


# CF vector example

$$CF = (N, \overrightarrow{LS}, SS)$$

$$\overrightarrow{LS} = \sum_{i=1}^N \vec{X}_i$$

$$SS = \sum_{i=1}^N \vec{X}_i^2 = \sum_{i=1}^N \langle X_i, X_i \rangle$$

(3,4)  
(4,5)  
(5,5)  
(3,6)  
(4,7)  
(3,8)



(5,1)  
(7,3)  
(7,4)  
(8,5)

$CF_1 = (6, (22,35), 299)$

$CF_2 = (4, (27,13), 238)$

## CF vector properties 1/4

- The CF vector is not only **efficient**, as it compresses the input dataset, but also **accurate**, as it is sufficient to compute several measures we need for clustering.

$$CF = (N, \overrightarrow{LS}, SS)$$

$$\overrightarrow{LS} = \sum_{i=1}^N \vec{X}_i$$

$$SS = \sum_{i=1}^N \vec{X}_i^2 = \sum_{i=1}^N \langle X_i, X_i \rangle$$

- the **centroid** of C: 
$$\vec{X}_0 = \frac{\sum_{i=1}^N \vec{X}_i}{N} \implies \frac{\overrightarrow{LS}}{N}$$

- the **radius** of C (avg distance from cluster members to the centroid):

$$R = \left( \frac{\sum_{i=1}^N (\vec{X}_i - \vec{X}_0)^2}{N} \right)^{\frac{1}{2}} \implies \sqrt{\frac{SS}{N} - \left( \frac{LS}{N} \right)^2}$$

**Homework:**  
**Prove it!**

- the **diameter** of C (avg pairwise distance within a cluster)

$$D = \left( \frac{\sum_{i=1}^N \sum_{j=1}^N (\vec{X}_i - \vec{X}_j)^2}{N(N-1)} \right)^{\frac{1}{2}} \implies$$

**Homework: express it in  
terms of CF statistics!**




## CF vector properties 2/4

- **CF additivity property:** Let two disjoint clusters  $C_1$  and  $C_2$ . The CF vector of the cluster that is formed by **merging** the two disjoint clusters, is:

$$CF(C_1 \cup C_2) = CF(C_1) + CF(C_2) = (N_1 + N_2, LS_1 + LS_2, QS_1 + QS_2)$$

- What is the CF of the marked (in red) cluster? How is it related to the CF1, CF2?



## CF vector properties 3/4

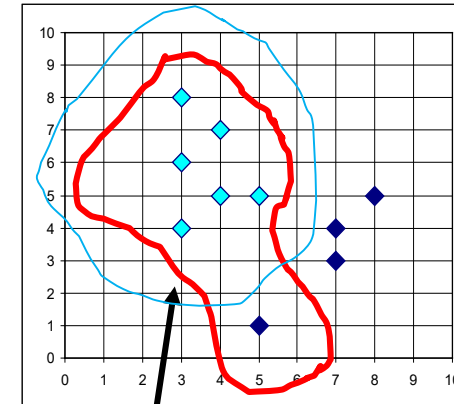
- **CF incremental property:** The updated CF of a cluster  $C_1$  after the **addition** of a new point  $p$ , is:

$$\text{CFT}(C_1 \cup p) = \text{CFT}(C_1) + p$$

- What is the CFP of the marked (in red) set?
- How is it related to CF1?

(3,4)  
(4,5)  
(5,5)  
(3,6)  
(4,7)  
(3,8)

(5,1)  
(7,3)  
(7,4)  
(8,5)



$CF_1 = (6, (22,35), 299)$

# CF vector properties 4/4

- Cluster feature (CF)
  - ❑ A summary of the statistics of the points in a cluster  $C$
  - ❑ Utilizes storage efficiently
  - ❑ Keeps sufficient statistics for clustering
  - ❑ Allows for easy **merge** of clusters, based on the **additivity property**
  - ❑ Allows for easy **addition of new points**, based on the **incremental property**

# Outline

- Data stream clustering basics
- Summarization
- Overview of stream clustering methods
- Partitioning methods
- Density-based methods
- Grid-based methods
- Data stream clustering: evaluation aspects
- Things you should know from this lecture & reading material


# A (non-complete) taxonomy of stream clustering approaches (& representative methods)

	Batch/Static clustering	Dynamic/Stream clustering
Partitioning methods	<ul style="list-style-type: none"> <li>• k-Means</li> <li>• k-Medoids</li> </ul>	<ul style="list-style-type: none"> <li>• <b>Leader</b></li> <li>• Simple single pass k-Means</li> <li>• <b>STREAM k-Means</b> [O'Callaghan et al 2002]</li> <li>• <b>CluStream</b> [Aggrawal et al 2003]</li> </ul>
Density-based methods	<ul style="list-style-type: none"> <li>• DBSCAN</li> <li>• OPTICS</li> </ul>	<ul style="list-style-type: none"> <li>• <b>DenStream</b> [Cao et al 2006]</li> <li>• incDBSCAN *</li> <li>• incOPTICS *</li> </ul>
Grid-based methods	<ul style="list-style-type: none"> <li>• STING</li> </ul>	<ul style="list-style-type: none"> <li>• <b>DStream</b> [Chen &amp; Tu 2007]</li> </ul>

(\*) These methods require access to the raw data (this access might be limited though)

# An example of an incremental clustering method - incDBSCAN

- Goal of **incremental methods**: To update the old clustering based on the new data (point  $p$ ), without reclustering the data from scratch.
  - Access to raw data is possible but unnecessary access should be avoided
  - It refers not only to adding a new point, but also to the removal of existing points
- incDBSCAN[Ester et al, 1998] exploits the locality of information in density-based clustering and reorganizes the information only locally (as required)
- In our example:
  - Only the affected cluster is re-organized, not everything is reclustered from scratch
  - Requires (limited) access to raw data (to the affected highlighted cluster in our example)



**Figure 3:** : Affected objects in a sample database

# Incremental clustering methods vs stream clustering methods

- Incremental methods require random access to the raw data to update the old clustering based on new instances.
  - They typically result in exact solutions
- Stream clustering methods do not assume random access to the data
  - The results are typically approximate
- Incremental methods might be appropriate for dynamic data arriving at a low rate
  - For potentially infinite streams, however they are not appropriate, and therefore new solutions are needed that can deal with the amount and complexity of the data


# Outline

- Data stream clustering basics
- Summarization
- Overview of stream clustering methods
- Partitioning methods
- Density-based methods
- Grid-based methods
- Data stream clustering: evaluation aspects
- Things you should know from this lecture & reading material

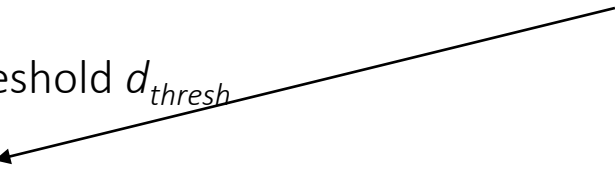


# Partitioning methods

- Goal: Construct a partition of a set of objects into  $k$  clusters so that some clustering criterion is optimized
  - e.g.  $k$ -Means,  $k$ -Medoids
- Two types of methods:
  - **Adaptive methods:**
    - Leader (Spath 1980)
    - Simple single pass  $k$ -Means (Farnstrom et al, 2000)
    - STREAM  $k$ -Means [OCaEtAl02]
  - **Online summarization - offline clustering methods:**
    - CluStream [Aggrwal et al, 2003]




# Adaptive methods: Leader (Spath 1980)

- The **simplest single-pass partitioning** algorithm
  - Whenever a new instance  $p$  arrives from the stream
    - Find its closest cluster (**leader**),  $c_{clos}$
    - Assign  $p$  to  $c_{clos}$  if their distance is below the threshold  $d_{thresh}$
    - Otherwise, create a new cluster (leader) with  $p$
  - Discussion
    - + 1-pass and fast algorithm
    - + No prior information on the number of clusters
    - The number of clusters is not controllable
    - It depends on the order of the examples
    - It depends on a correct guess of  $d_{thresh}$
- A cluster is “defined” by its first point
- 

# Adaptive methods: STREAM k-Means [O'Callaghan et al 2002]


- An extension of k-Means for streams
  - The iterative process of static  $k$ -Means cannot be applied to streams
  - Idea: Use a **buffer** that fits in memory and apply  $k$ -Means **locally** in the buffer
- Stream is processed in chunks  $X_1, X_2 \dots X_i \dots$ , each fitting in memory
  - For the current chunk  $X_i$ 
    - Apply  $k$ -Means locally on  $X_i$  (retain only the  $k$  cluster centers from  $X_i$ )
    - $X'$ : the cluster centers seen thus far over the stream ( $\# i * k$  centers)
      - Each center is treated as a point, weighted with the number of points it compresses
    - Apply  $k$ -Means on  $X'$  to obtain the current clustering result



$X_1$

...


...



$X_i$


# Adaptive methods: STREAM k-Means [O'Callaghan et al 2002]

- For each batch




# Adaptive methods: STREAM k-Means [O'Callaghan et al 2002]

- For each batch



# Adaptive methods: STREAM k-Means [O'Callaghan et al 2002]

- An [example of Stream k-Means in SPARK](#)



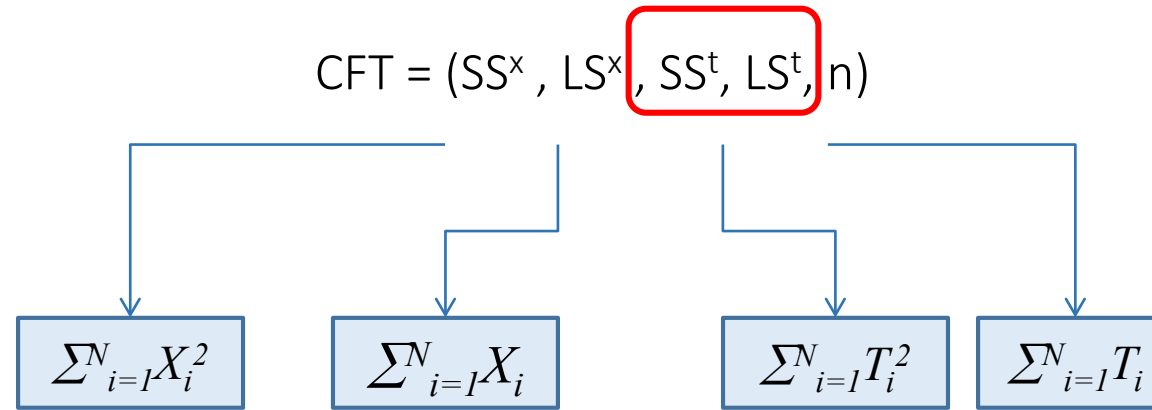
- Discussion
  - + The number of clusters is controllable ( $k$  on each batch)
  - + Good results on each batch (via typical  $k$ -Means)
  - The number of clusters is fixed over the stream ( $k$ )
  - Only one clustering model is reported at each time point

# Online-Offline methods: CluStream [Aggrawal et al, 2003]

- The stream clustering process is separated into two components:
  - an **online micro-cluster component**, that summarizes the stream locally as new data arrive over time
    - Micro-clusters are stored in disk at snapshots in time that follow a pyramidal time frame.
  - an **offline macro-cluster component**, that clusters these summaries into global clusters
    - Clustering is performed upon summaries instead of raw data

# CluStream: the micro-cluster summary structure

- The microcluster summaries are extensions of the cluster feature vector (CF) summary of BIRCH
- The micro-cluster summary for a set of  $d$ -dimensional points  $(X_1, X_2, \dots, X_n)$  arriving at time points  $T_1, T_2, \dots, T_n$  is defined as:




LS stands for Linear Sum  
SS stands for Square Sum



# CluStream: the micro-cluster summary structure

- The micro-cluster summary for a set of d-dimensional points  $(X_1, X_2, \dots, X_n)$  arriving at time points  $T_1, T_2, \dots, T_n$  is defined as:



LS stands for Linear Sum  
SS stands for Square Sum

- Using the summaries, we can easily calculate basic measures to characterize a cluster:

- Center:  $\vec{X}_0 = \frac{\sum_{i=1}^N \vec{X}_i}{N} \Rightarrow \frac{\vec{LS}}{N}$


- Radius:  $R = \left( \frac{\sum_{i=1}^N (\vec{X}_i - \vec{X}_0)^2}{N} \right)^{\frac{1}{2}} \Rightarrow \sqrt{\frac{SS}{n} - \left( \frac{LS}{n} \right)^2}$

....

- Similarly information on the cluster recency can be derived, e.g., avg cluster timestamp

# CluStream: the micro-cluster summary structure

- In other words, we summarize the stream via micro-clusters
  - ▣ Each microcluster is represented through its CFT summary



# CluStream: the micro-cluster summary structure

- Micro-clusters have very appealing properties for streams

- **Incrementality:**  $CFT(C_1 \cup p) = CFT(C_1) + p$

- we can easily add new points to a microcluster

- **Additivity:**  $CFT(C_1 \cup C_2) = CFT(C_1) + CFT(C_2)$

- we can easily merge two microclusters

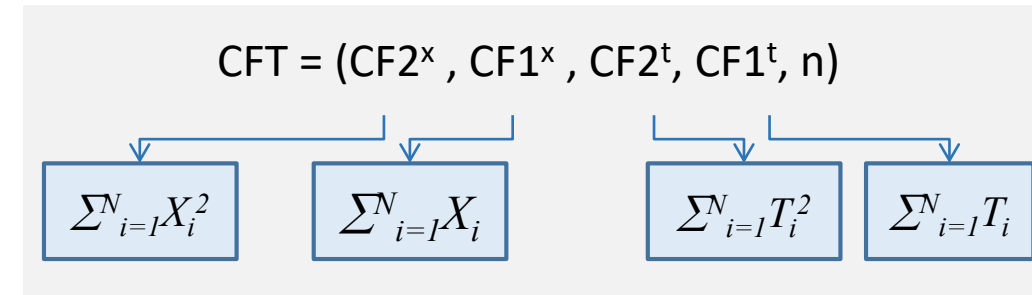
- **Subtractivity:**  $CFT(C_1 - C_2) = CFT(C_1) - CFT(C_2)$ ,  $C_1 \supseteq C_2$

- we can remove the effect of an old microcluster

- Recall the 2-directional learning in streams

- Incrementality helps us to incorporate new information in the clustering model

- Subtractivity helps us to remove outdated information from the clustering model

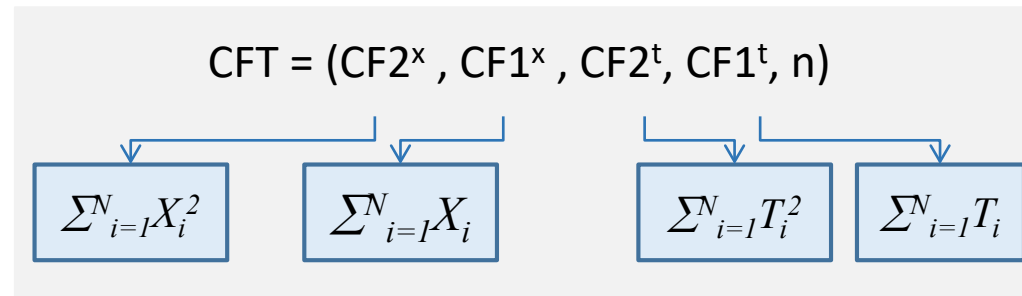


# CluStream algorithm: overview


- Input:
  - The stream
  - $q$ : #micro-clusters to be maintained over time (fixed)
  - $t$ : radius factor
- 4 steps
  - **Initialization**: How we build the initial set of microclusters?
  - **Online** micro-cluster maintenance: How do we add new points from the stream?
  - **Periodic** storage: Decide when to store snapshots of micro-clusters on disk?
  - **Offline** macro-clustering: How to derive the final clusters?

# CluStream: Initialization step

- Initialization: How we build the initial set of microclusters?
  - ❑ Done using an offline process in the beginning of the stream
  - ❑ Wait for the first *InitNumber* points to arrive
  - ❑ Apply a standard *k*-Means algorithm with  $k=q$  to create  $q$  clusters
  - ❑ For each discovered cluster, assign it a unique ID and create its micro-cluster summary.




- How should we set  $q$ ?
  - ❑ much larger than the natural number of clusters
  - ❑ much smaller than the total number of points arrived




# CluStream: Online step - How do we add new points from the stream?

- A fixed number of  $q$  micro-clusters is maintained over time
- Whenever a new point  $p$  arrives from the stream
  - Compute distance between  $p$  and each of the  $q$  maintained micro-cluster centroids
  - $clu \leftarrow$  the closest micro-cluster to  $p$
  - Find the max boundary of  $clu$ 
    - It is defined as a factor of  $t$  of  $clu$  radius
  - If  $p$  falls within the maximum boundary of  $clu$ 
    - $p$  is **absorbed** by  $clu$
    - Update  $clu$  statistics (**incremental property of microclusters**)
  - Else, create a **new** micro-cluster with  $p$ , assign it a new ID, initialize its statistics
    - To keep the total number of micro-clusters fixed (i.e.,  $q$ ):
      - **Delete** the most obsolete micro-cluster or
        - If it is safe based on its time statistics
      - **Merge** the two closest ones (**Additivity property of microclusters**)
        - When two micro-clusters are merged, a **list of ids** is created. This way, we can identify the component micro-clusters that comprise a micro-cluster.




The green point can be absorbed by the summary.




The green point cannot be absorbed by the summary.

# CluStream: Periodic micro-cluster storage

- Micro-clusters are stored as snapshots in time following the pyramidal pattern framework



# CluStream: Offline step




- The offline step is applied on demand. The user specifies the **time horizon  $h$  for clustering**, e.g.,  $h=10$ , so from  $S_T - S_{T-10}$  where  $T$  is the current timepoint.
  - Different clusterings are possible, if the horizon of clustering changes → allows the user can explore the history of the stream
- User input: time horizon  $h$ , # macro-clusters  $k$  to be detected, current time  $T$
- Output: the clusters in  $(T-h, T)$

2 steps

- Step 1: Find the **active micro-clusters** during  $h$
- Step 2: Apply k-Means over the active micro-clusters in  $h$  to derive the  $k$  macro-clusters




# CluStream: Offline step – step 1



- **Step 1:** Find the active micro-clusters during horizon  $h$ :
- We exploit the subtractivity property to find the active micro-clusters during  $h$
- Suppose current time is  $T$ . Let  $S(T)$  be the set of micro-clusters at  $T$ .
- Find the stored snapshot which occurs just before time  $T-h$ . Let  $S(T-h')$  be the set of micro-clusters.
- For each micro-cluster in the current set  $S(T)$ , we find the list of its component micro-cluster ids. For each of the list of ids, find the corresponding micro-clusters in  $S(T-h')$ .
- Subtract the CF vectors for the corresponding micro-clusters in  $S(T-h')$  (**subtractivity property of microclusters**)
- This ensures that the micro-clusters created before the user-specified horizon do not dominate the result of clustering process

# CluStream: Offline step – step 1 - example

- **Example:** if we have a merged cluster with id list (C1,C2,C3) in  $S(T)$  and a cluster with ID C1 in  $S(T-h')$ , then we should remove C1 as it was created before the user horizon.
  - To this end, we can use the subtractivity property:  $CFT(C1,C2,C3) - CFT(C1) = CFT(C2,C3)$
  - The result is the active IDs during h



## CluStream: Offline step – step 2

- **Step 2:** Apply k-Means over the active micro-clusters in  $h$  to derive the  $k$  macro-clusters
  - Initialization: centers are not picked up randomly, rather sampled with probability proportional to the number of points in a given micro-cluster
  - Distance is the centroid distance
  - New centers are defined as the weighted centroids of the micro-clusters in that micro-cluster partition

# CluStream algorithm: overview

Input:       The stream,

$q$ : #micro-clusters to be maintained over time (fixed)

$t$ : radius factor

- **Initialize**: How we build the initial set of microclusters
  - ▣ Apply  $k$ -Means, with  $k=q$ , over *initPoints*. Built a summary for each cluster.
- **Online micro-cluster maintenance**: How do we add new points from the stream?
  - ▣ Does a new point fit into some existing summary?
    - If yes, just update the summary
    - If no, create a new summary. But because of fixed  $q$ , you have to reduce #microclusters by one
- **Periodic storage of micro-clusters snapshots into disk**
  - ▣ At different levels of granularity depending upon their recency
- **Offline (on demand) macro-clustering**
  - ▣ Input: A user defined time horizon  $h$  and number of macro-clusters  $k$  to be detected
  - ▣ Locate the valid micro-clusters during  $h$
  - ▣ Apply  $k$ -Means upon these micro-clusters →  $k$  macro-clusters

# CluStream: discussion

- + One pass over the raw data
- + Views the stream as a changing process over time, rather than clustering the whole stream
- + Provides flexibility to an analyst in a real-time and changing environment
- + Can characterize clusters over different time horizons in a changing environment
- Fixed number of micro-clusters maintained over time
- Sensitive to outliers/ noise
  - We might delete a valid microcluster just because of an outlier point

# Homework (could also be a potential example topic)

## ■ Data stream clustering: ageing/window model



- We discussed different forms of ageing in a stream environment: landmark windows, sliding windows, damped window models, ...
- Questions:
  - i) Please explain the ageing schema in the CluStream algorithm.
  - ii) Can it be categorized as sliding or landmark or damped window model?
- Answer:
  - Landmark: Yes/No. Why?
  - Sliding: Yes/No. Why?
  - Damped: Yes/No. Why?

# Outline

- Data stream clustering basics
- Summarization
- Overview of stream clustering methods
- Partitioning methods
- Density-based methods
- Grid-based methods
- Data stream clustering: evaluation aspects
- Things you should know from this lecture & reading material

# Density-based methods

- Clusters as regions of high density surrounded by regions of low density (noise)
  - Density is measured **locally**, in the  $\epsilon$ -neighborhood of each point
    - e.g. DBSCAN, OPTICS
- Very appealing for streams
  - No assumption on the number of clusters
  - Discovering clusters of arbitrary shapes
  - Ability to handle outliers and noise
- But, they miss a clustering model (or it is too complicated)
  - Clusters are represented by all their points!!!!
- Solution: **Describe clusters as set of summaries**
  - DenStream [Cao et al 2006]






# DenStream [Cao et al 2006]

- The **online-offline rationale** is followed:
  - **Online summarization** as new data arrive over time
    - They distinguish between different types of summaries: Core, potential core and outlier micro-clusters
  - **Offline clustering** over the summaries to derive the final clusters
    - A modified version of DBSCAN over the summaries
- Data are subject to ageing according to the exponential ageing function (**damped window model**) – recall previous lectures on streams

$$f(o, t) = e^{-\lambda(t-t_o)}$$


- $t-t_o$  is the time from point occurrence
- $\lambda$  ( $\lambda > 0$ ) is the decay rate which determines the importance of historical data
- The higher the value of  $\lambda$ , the lower the importance of old data



# DenStream: summarizing the stream

Note that the microcluster is now a temporal object

- The **micro-cluster summary at time  $t$**  for a set of  $d$ -dimensional points  $(p_1, p_2, \dots, p_n)$  arriving at time points  $T_1, T_2, \dots, T_n$  is:



- Easy computation of basic measures, e.g., (see also discussion on micro-cluster summaries):

- Center:  $c = \frac{CF^1}{w}$

- Radius:  $r = \sqrt{\frac{CF^2}{w} - \left(\frac{CF^1}{w}\right)^2}$

- A micro-cluster summary  $c_p$  can be **maintained incrementally**

- If **a new point  $p$  is added** to  $c_p$ :

$$c_p = (CF^1+p, CF^2+p^2, w+1)$$

- If **no point is added** to  $c_p$  for time interval  $\delta t$ , decay should be applied:

$$c_p = (2^{-\lambda\delta t}*CF^1, 2^{-\lambda\delta t}*CF^2, 2^{-\lambda\delta t}*w)$$

# DenStream: core, potential core & outlier summaries

## User-defined parameters:

$\mu$ : the density threshold for core-microclusters (similar to DBSCAN)

$\epsilon$ : the radius threshold (similar to DBSCAN)

$\beta$ :  $\beta * \mu$  the density threshold for potential-core microclusters

### ■ Core (or dense) micro-clusters

- $(w \geq \mu) \ \& \ (r \leq \epsilon)$

### ■ But, in an evolving stream, the role of clusters and outliers often interchange:


- Should provide opportunity for the gradual growth of new clusters
- Should promptly get rid of the outliers

### ■ Potential core micro-clusters

- $(w \geq \beta * \mu) \ \& \ (r \leq \epsilon), \ 0 < \beta \leq 1$

### ■ Outlier micro-clusters

- $(w < \beta * \mu) \ \& \ (r \leq \epsilon), \ 0 < \beta \leq 1$




# DenStream: the algorithm

- ❑ The algorithm consists of 4 main components:
  - **Initialize**: apply DBSCAN over initPoints → the core points are the p-micro-clusters
  - **Online step**: Maintain the micro-clusters as new points arrive from the stream
    - ❑ 2 lists of **p-micro-clusters** and **o-micro-clusters** are maintained over time
  - **Periodic micro-cluster maintenance** due to data ageing
  - **Offline macro-clustering**: upon user request, extract the final clusters

# DenStream: online step

- 2 lists of p-micro-clusters and o-micro-clusters are maintained over time
- When a new point  $d$  arrives
  - Find its closest p-micro-cluster  $pclu$ 
    - If the updated radius of  $pclu \leq \epsilon$ , merge  $d$  to  $pclu$
  - otherwise find its closest o-micro-cluster  $oclu$ 
    - If the updated radius of  $oclu \leq \epsilon$ , merge  $d$  to  $oclu$
    - Check if  $oclu$  can be upgraded to a p-micro-cluster (if now  $w \geq \beta * \mu$ )
  - o.w., create a new o-micro-cluster with  $d$  (keep also the creation time  $t_o$  for the microcluster)



# DenStream: periodic microcluster maintenance 1/2

- DenStream maintains 2 separate memories: i) for p-micro-clusters and ii) o-micro-clusters
- How to ensure that the memory is bounded?
  - General principle: delete outdated information
- For **p-micro-clusters memory**, the idea is to delete p-micro-clusters that turn into o-micro-clusters.
  - This means, checking the weight of each p-micro-cluster and delete those with weight  $< \beta^*\mu$
- The question is **how often should we check the weight** (so, efficiency).
  - Recall that  $\beta^*\mu$  is the minimum weight of a p-micro-cluster.
  - Therefore the minimum time for a p-micro-cluster to fade into an o-micro-cluster is given by:

$$2^{-\lambda T_p} \beta^*\mu = \beta^*\mu - 1 \quad \rightarrow \quad T_p = \lceil \frac{1}{\lambda} \log(\frac{\beta^*\mu}{\beta^*\mu - 1}) \rceil,$$

- So, we need to check for deletion of p-micro-clusters every  $T_p$  time periods
- This checking strategy also ensures that **the p-microclusters memory is bounded** by  $W / \beta^*\mu$  where the constant  $W$  is the overall weight of the datastream (proof omitted, see paper)

## DenStream: periodic microcluster maintenance 2/2



- DenStream maintains two separate memories: i) for p-micro-clusters and ii) o-micro-clusters
- How to ensure that the memory is bounded? General principle: delete outdated information
- For **o-micro-clusters memory**, the problem is that their number might continuously grow. We need to keep them as an o-micro-cluster might be the beginning of a p-micro-cluster. But we cannot keep them forever.
- So the idea is to delete those o-micro-clusters that are **not promising anymore**.
- How **promising** a microcluster is depends on its **current weight  $w$**  vs its **expected weight  $\xi$** :

$$\xi(t_c, t_o) = \frac{2^{-\lambda(t_c - t_o + T_p)} - 1}{2^{-\lambda T_p} - 1}$$

- Intuitively, the longer a microcluster exists (larger  $(t_c - t_o)$ ), the higher its weight is expected to be.
- If  $w < \xi$ , the o-micro-cluster might not grow into a p-micro-cluster and can be safely deleted.
- The authors prove that **the number of o-micro-clusters is bounded** (proof omitted, see paper)

# DenStream: offline step

- Upon request, apply a **variant of DBSCAN** over the set of online maintained p-micro-clusters
  - Each p-micro-cluster  $c_p$  is treated as a virtual point located at the center of  $c_p$  with weight  $w$ .
- Core-micro-clusters (redefined)
- Directly density reachable (redefined)
  - $c_p$  is directly density reachable from  $c_q$  if:
    - $c_q$  is a c-micro-cluster and
    - $\text{dist}(c_p, c_q) \leq 2\epsilon$  (i.e. they are tangent or intersecting)
- Density reachable (redefined)
  - A p-micro-cluster  $c_p$  is density reachable from a c-micro-cluster  $c_q$  if there is a chain of c-micro-clusters  $c_{p1}=c_q, c_{p2}, \dots, c_{pn}=c_p$ .
- Density connected (redefined)
- A cluster is a maximum set of density connected points





# DenStream: overview


- + DenStream clusters large evolving data streams
  - + Discover clusters of arbitrary shapes, following the density-based paradigm
  - + No assumption on the number of clusters
  - + Noise/ outlier handling
- 
- The choice of the parameters  $\epsilon$ ,  $\beta$ ,  $\mu$
  - Constant parameters over time, what about clusters with different density

# Outline

- Data stream clustering basics
- Summarization
- Overview of stream clustering methods
- Partitioning methods
- Density-based methods
- Grid-based methods
- Data stream clustering: evaluation aspects
- Things you should know from this lecture & reading material


# Grid based methods

- A grid structure is used to capture the density of the dataset.
  - A cluster is a set of connected dense cells
  - e.g. STING
- Appealing features for streams
  - No assumption on the number of clusters
  - Discovering clusters of arbitrary shapes
  - Ability to handle outliers
- In case of streams
  - The grid cells “constitute” the summary structure
  - Update the grid structure as the stream proceeds
  - DStream [Chen & Tu 2007]



## DStream [Chen & Tu 2007]

- Resembles the online-offline rationale of CluStream/DenStream but there is no real offline part, rather a final clustering structure is maintained online.
  - Online mapping of the new data into the grid (so summarization)
  - Periodic final clustering maintenance




# DStream: Summarizing the stream into the grid

Note that the cell-summary is a temporal object

- Data ageing (**damped window model**):
  - $D(x,t) = \lambda^{t-t_c}$ ,  $t_c$  is the arrival time for point  $x$ ,  $t$  is the current timepoint
  - $\lambda$  in  $(0,1)$  is the *decay factor*

- The **density of a grid cell  $g$  at time  $t$** : 
$$D(g,t) = \sum_{x \in E(g,t)} D(x,t)$$

- The characteristic vector of a grid cell  $g$  is defined as:



- The grid density can be updated incrementally

$$D(g, t_n) = \lambda^{t_n - t_l} D(g, t_l) + 1$$


$t_n$ : the new record arrival time;  $t_l$ : the last record arrival

# DStream: Dense, Sparse and Transitional grid cells

- The density of a grid is constantly changing over time.


- Dense grid cells

$$D(g, t) \geq \frac{C_m}{N(1 - \lambda)} = D_m \quad C_m > 1$$




- Transitional grid cells

$$D(g, t) \leq \frac{C_l}{N(1 - \lambda)} = D_l \quad 0 < C_l < 1$$

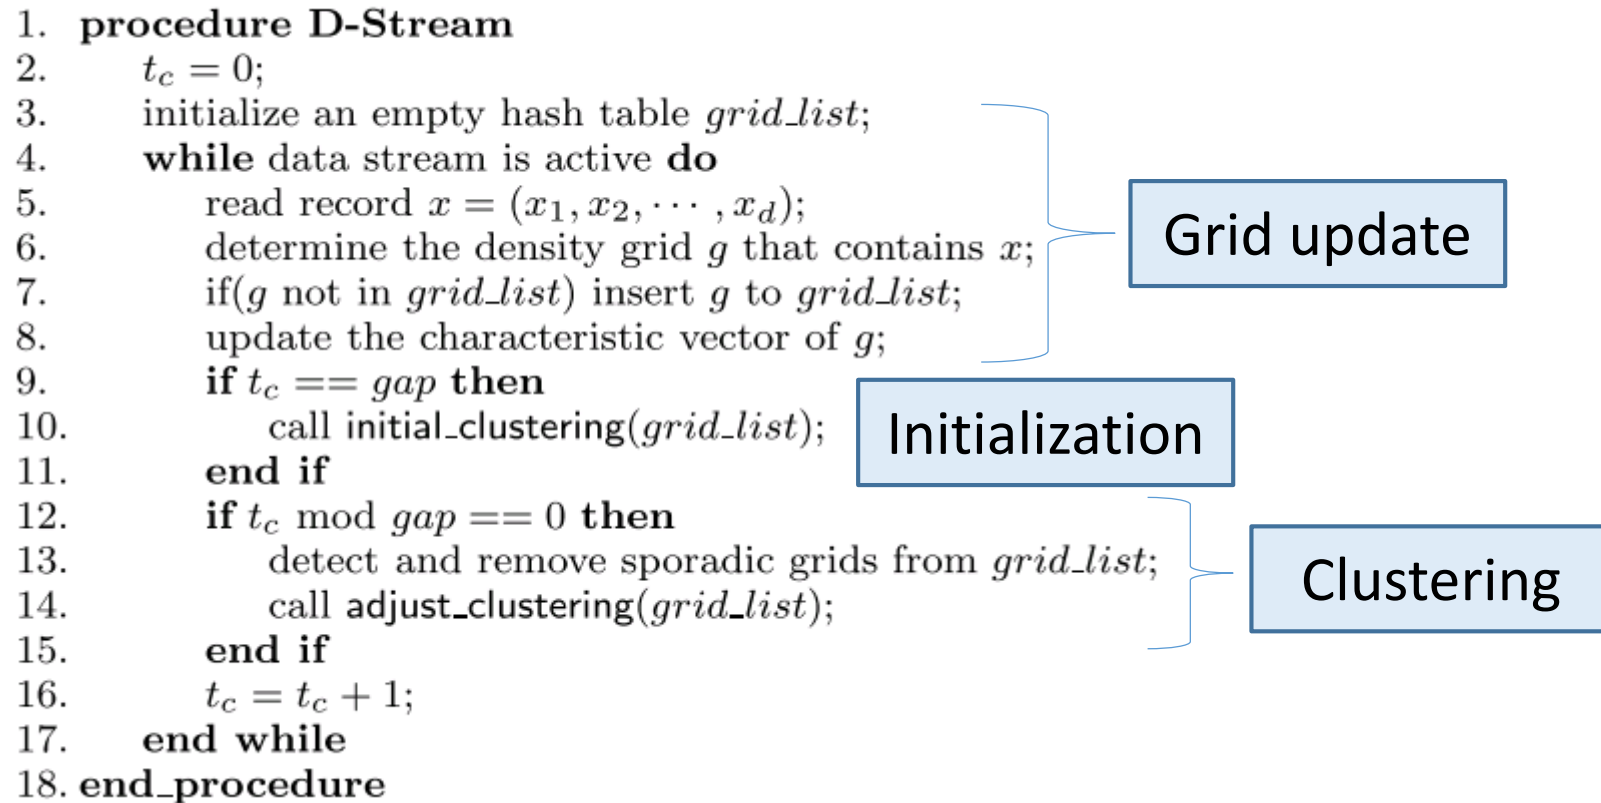


- Sparse grid cells

$$\frac{C_l}{N(1 - \lambda)} \leq D(g, t) \leq \frac{C_m}{N(1 - \lambda)}$$



# DStream: the algorithm



# DStream: overview

- + DStream clusters large evolving data stream
  - + It can discover clusters of arbitrary shapes
  - + No assumption on the number of clusters
  - + Distinguishes noise and outliers
  - + The grid provides a level of abstraction over the data
- 
- The choice of the grid parameters
  - Fixed grid parameters



# Outline

- Data stream clustering basics
- Summarization
- Overview of stream clustering methods
- Partitioning methods
- Density-based methods
- Grid-based methods
- Data stream clustering: evaluation aspects
- Things you should know from this lecture & reading material

# Stream clustering evaluation 1/3

- Similar to what we discussed for batch clustering, but in streams we are interested also in the overtime performance monitoring of the clustering
- Two clustering quality categories (similarly to batch evaluation, see corresponding lectures)
  - **Internal measures of similarity**, e.g., SSE measuring the goodness of a clustering structure
  - **External measures of similarity**, e.g., entropy, measuring the extent to which cluster labels match externally supplied class labels.
- The difference is that those measures are evaluated over a **future horizon**




Figure 4: Quality comparison (Charitable Donation dataset, horizon=16, stream\_speed=200)

## Stream clustering evaluation 2/3

- Except for the quality, other important aspects
  - ▣ **Time** (how fast the points are processed)




Figure 14: Execution time vs. length of stream(Network Intrusion data set)

# Stream clustering evaluation 3/3

- Except for the quality, other important aspects

- **Memory**

- For example, for methods that not assume a fixed number of summaries, plotting this number over time is very informative about the underlying population distribution




Figure 5: Number of (micro)clusters (Network Intrusion dataset, window size  $w=1000$ )

# Outline

- Data stream clustering basics
- Summarization
- Overview of stream clustering methods
- Partitioning methods
- Density-based methods
- Grid-based methods
- Data stream clustering: evaluation aspects
- Things you should know from this lecture & reading material

# Stream clustering overview

- A very important task given the availability of streams nowadays
- Stream clustering algorithm maintain a valid clustering of the evolving stream population over time
- Two generic approaches
  - Online maintenance of a final clustering model
  - Online summarization of the stream and offline clustering
    - Summaries!
- Different window models
- Handling outliers (or, potential future clusters) is very important
- Specialized approaches for text streams, high-dimensional streams.
- Evaluation is not straightforward

# Hands on experience



- Familiarize yourself with popular frameworks for stream learning
  - **MOA**: the most popular open source framework for data stream mining  
<https://moa.cms.waikato.ac.nz/> (in Java)
    - “[Machine Learning for Data Streams with Practical Examples in MOA](#)” book
  - Scikit-multiflow “A machine learning package for streaming data in Python”
    - <https://scikit-multiflow.github.io/>

■



## Related Open Source Software

- [RIVER](#), a new framework for stream mining in Python.
- [streamDM for Spark Streaming](#), a new framework for Spark.
- [Apache SAMOA](#), a new framework for distributed stream mining, can be easily used with Apache Flink, Apache Storm, S4, or Samza.
- [streamDM C++](#), a framework in C++ for data stream mining.
- [ADAMS](#), a novel, flexible workflow engine, is the perfect tool for maintaining MOA real-world, complex knowledge workflows.
- [The MEKA project](#) provides an open source implementation of methods for multi-label classification and evaluation.

Source: <https://moa.cms.waikato.ac.nz/>

# Reading material

- Book: Knowledge discovery from data streams, J. Gama
- Book: Data streams – Models and Algorithms, C. Aggrawal
- [Zhang et al 1996] [BIRCH: an efficient data clustering method for very large databases](#)
- [Breuning et al 2001] [Data Bubbles: Quality Preserving Performance Boosting for Hierarchical Clustering](#)
- [Aggrawal et al 2003] [A framework for clustering evolving data streams](#)
- [Cao et al 2006] [Density-Based Clustering over an Evolving Data Stream with Noise](#)
- [Chen & Tu 2007] [Density-Based Clustering for Real-Time Stream Data](#)
- F. Farnstrom, J. Lewis, C. Elkan: Scalability for clustering algorithms revisited. ACM SIGKDD Explorations Newsletter 57, 2000.
- S. Guha, A. Meyerson, N. Mishra, R. Motwani, L. O' Callaghan: Clustering data streams: Theory and practice. Data Mining and Knowledge Discovery 15(3):515–528, 2003.
- [O'Callaghan et al 2002] L. O'Callaghan, N. Mishra, A. Meyerson, S. Guha, R. Motwani: Streaming-Data Algorithms for Quality Clustering. ICDE, 2002.
- [Ester et al 1998] Ester et al, Incremental Clustering for Mining in a Data Warehousing Environment, VLDB 1998.





Thank you

Questions/Feedback/Wishes?

# Acknowledgements

- The slides are based on
  - ▣ *DM2 lecture@LUH(@Eirini Ntoutsi), KDD2/SS16 lecture@LMU Munich (@Eirini Ntoutsi, Matthias Schubert, Arthur Zimek)*